

# Scheduling & Orchestration

---

5th High Performance Container Workshop - ISC19

# Scope and Introduction

This segment focuses on **SCHEDULING & ORCHESTRATION** aspects.

We are going to touch on what was said before.



# Docker SWARM: Light-weight orchestrator

---

HPCW @ ISC2019

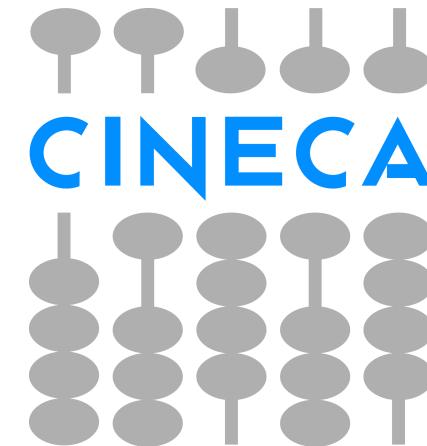
Abdulrahman Azab

Dept. of Research Computing

University of Oslo, Norway



# PRACE 6IP WP6.2.3: The deployment of containers and full virtualised tools into HPC infrastructures





# Docker UCP

**PRACE**

**Docker Enterprise Universal Control Plane v3.1.0-rc1**

**MANAGER NODES**

Status	Count	Errors	Pending
Ready	1	1	0
Warnings	0	0	Pending

**WORKER NODES**

Status	Count	Errors	Pending
Ready	2	2	0
Warnings	0	0	Pending

**LAST 6 HOURS**

**1 MANAGER NODE**

Max CPU 22.7% Max Memory 23.38% Max Used Disk 13.22%

**2 WORKER NODES**

Max CPU 3.55% Max Memory 2.78% Max Used Disk 10.44%

**SWARM**

- Services
  - Active: 0
  - Errors: 0
  - Updating: 0

**KUBERNETES**

- default
  - Pods
    - Running: 0
    - Errors: 0
    - Pending: 0
  - Controllers: None

**Docs**

- Kubernetes API Docs
- Live API

**Add Nodes**

Scale your swarm by adding additional worker nodes. Add manager nodes for high availability.

**Docker CLI**

Download a client bundle to create and manage services using the Docker CLI client.

**Manage Users & Teams**

Manage users and permissions by creating user accounts or integrating with an existing LDAP server.

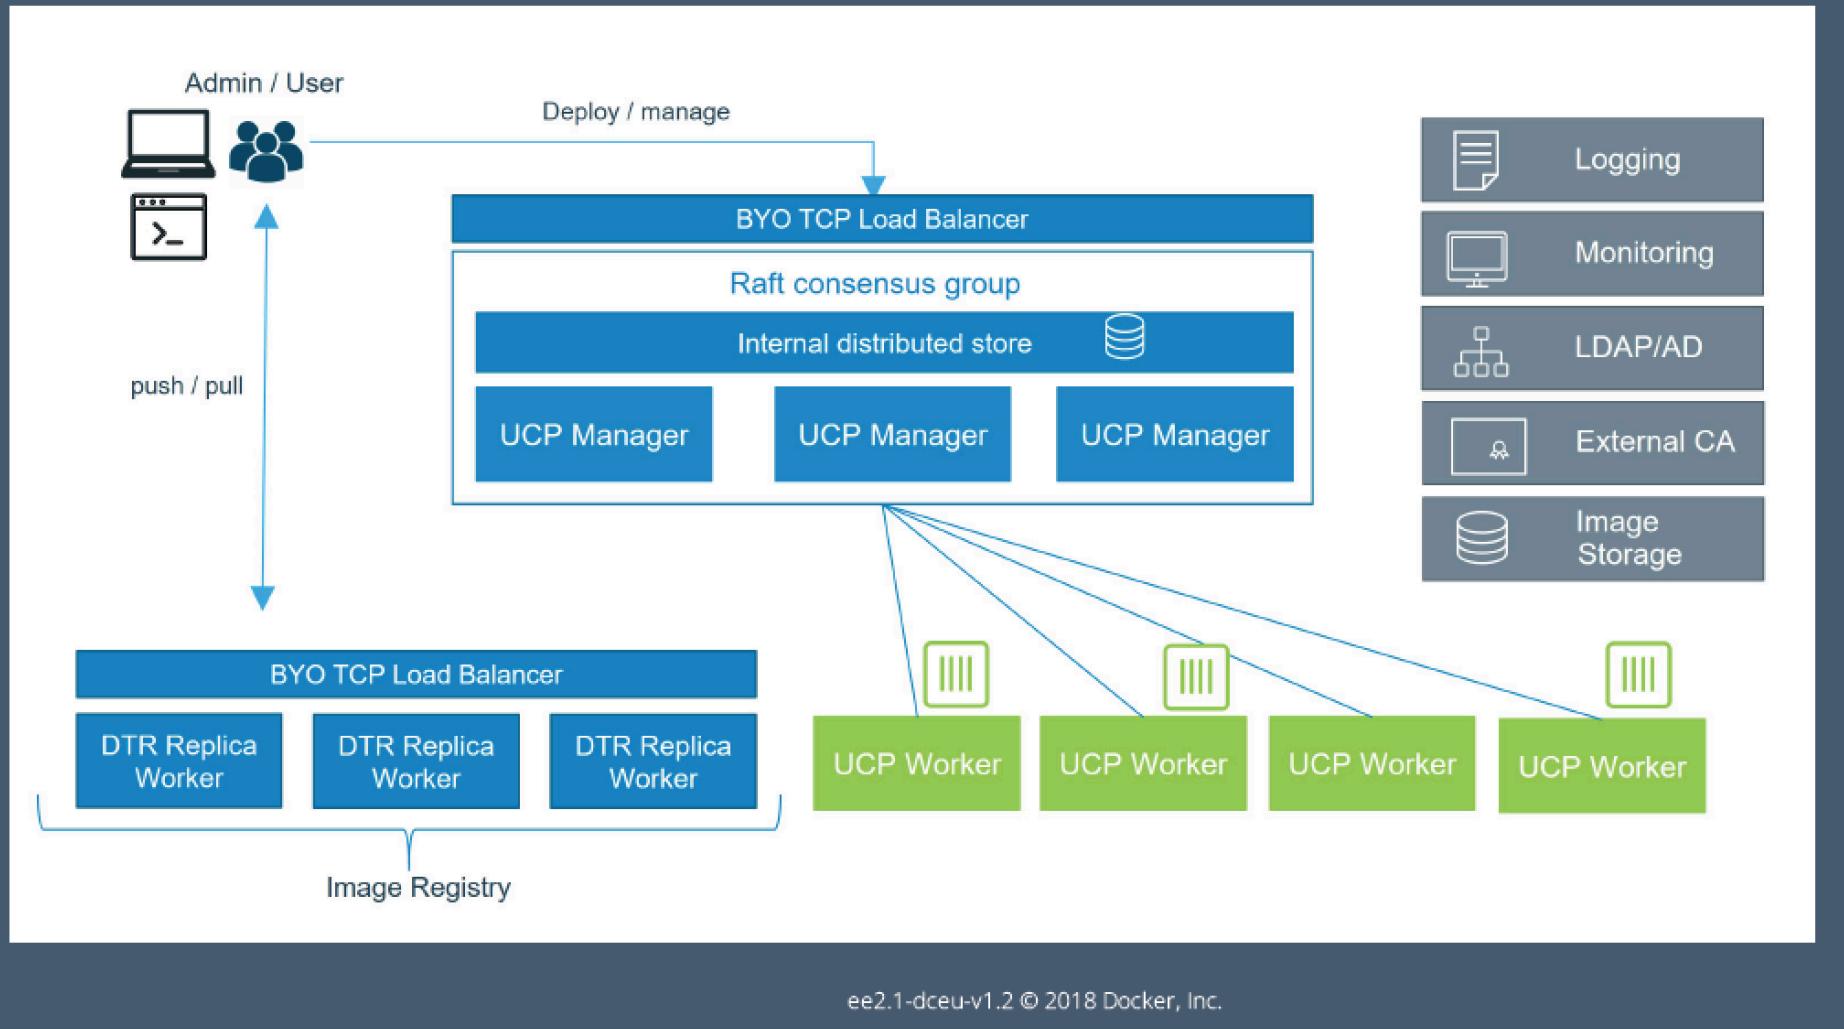
**Content Trust**

Configure UCP to only run services that use images signed by publishers you trust.

**Access Control**

Control who can access a set of resources by grouping those resources in collections.

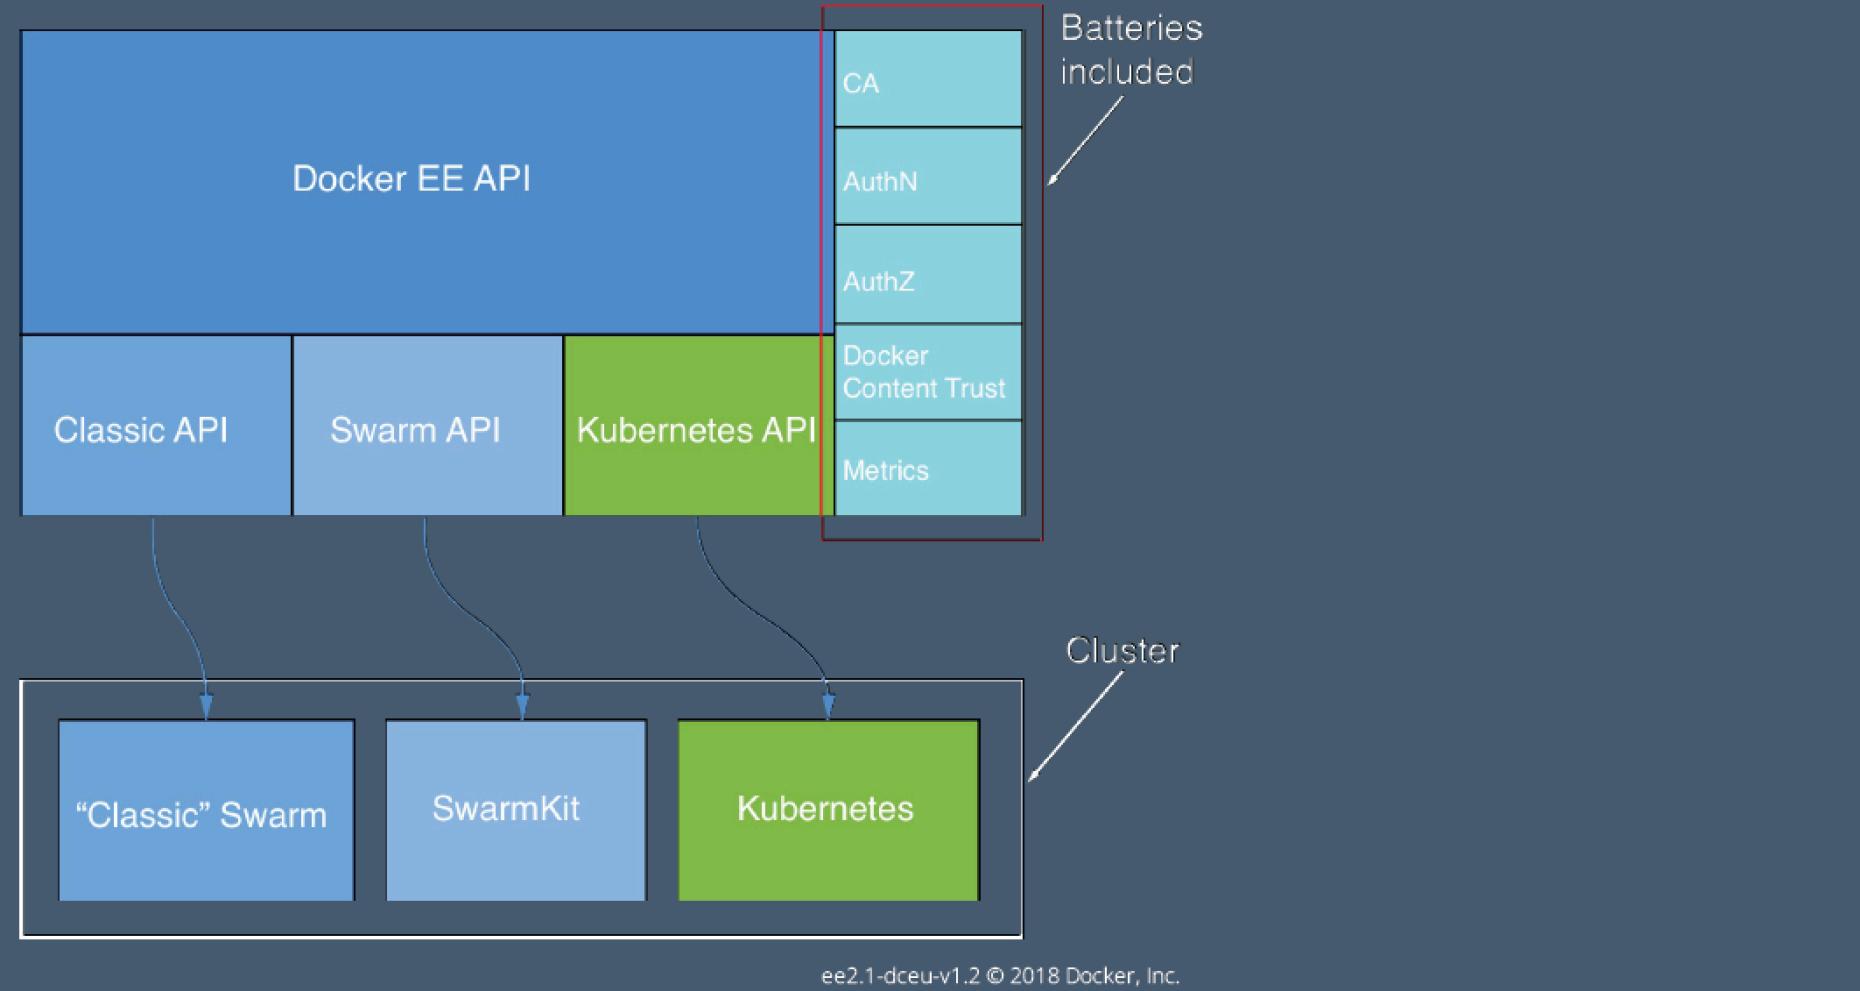
# DOCKER EE ARCHITECTURE





```
[centos@ucp-manager-0 ~]$ UCP_IP=<ucp-manager-0 IP>
[centos@ucp-manager-0 ~]$ UCP_FQDN=<ucp-manager-0 FQDN>
[centos@ucp-manager-0 ~]$ docker container run --rm -it --name
ucp \
-v /var/run/docker.sock:/var/run/docker.sock \
docker/ucp:3.1.0 install \
--admin-username admin \
--admin-password adminadmin \
--san ${UCP_IP} \
--san ${UCP_FQDN}
```

# CHOOSE YOUR ORCHESTRATOR





# Orchestration

# ORCHESTRATORS IN UCP

Managers



Docker Engine

swarmkit

ucp-kube-api-server

ucp-kube-controller-manager

ucp-kube-scheduler

Swarm Orch  
Kube Orch

Host OS

Workers

ucp-kube-proxy

ucp-kubelet

ucp-kube-proxy

ucp-kubelet

Docker Engine

swarmkit

Docker Engine

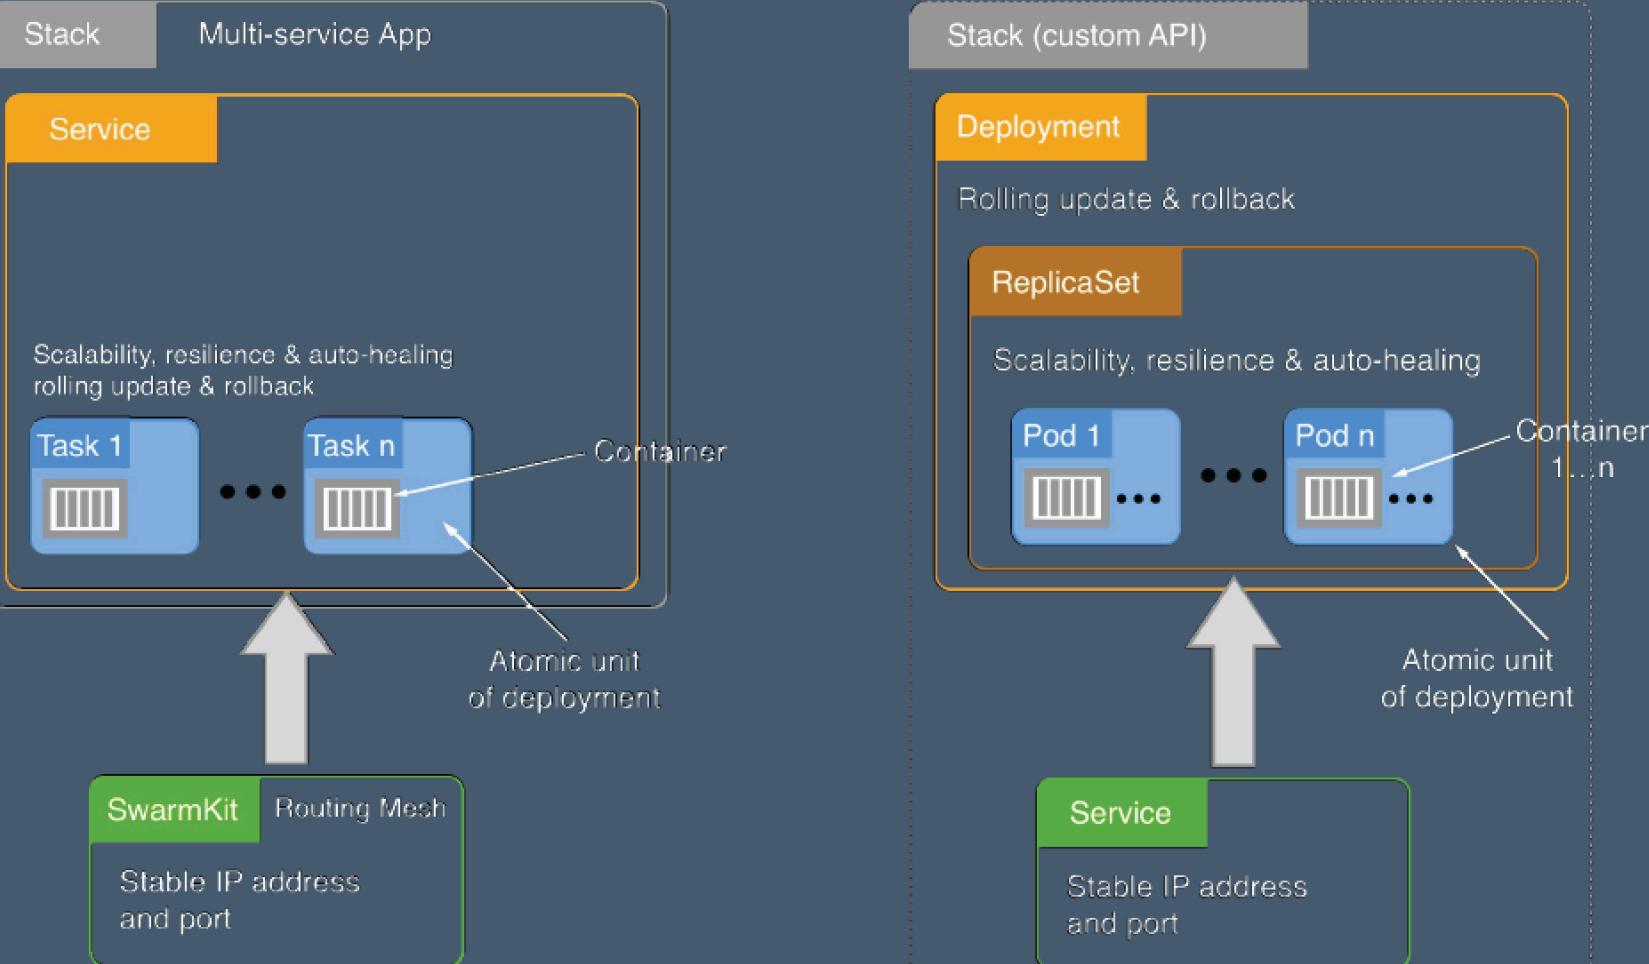
swarmkit

Host OS

Host OS

ee2.1-dceu-v1.2 © 2018 Docker, Inc.

# ORCHESTRATION COMPONENTS



ee2.1-dceu-v1.2 © 2018 Docker, Inc.



# Container Network Operations

# CHOOSING THE RIGHT NETWORKING TOOL

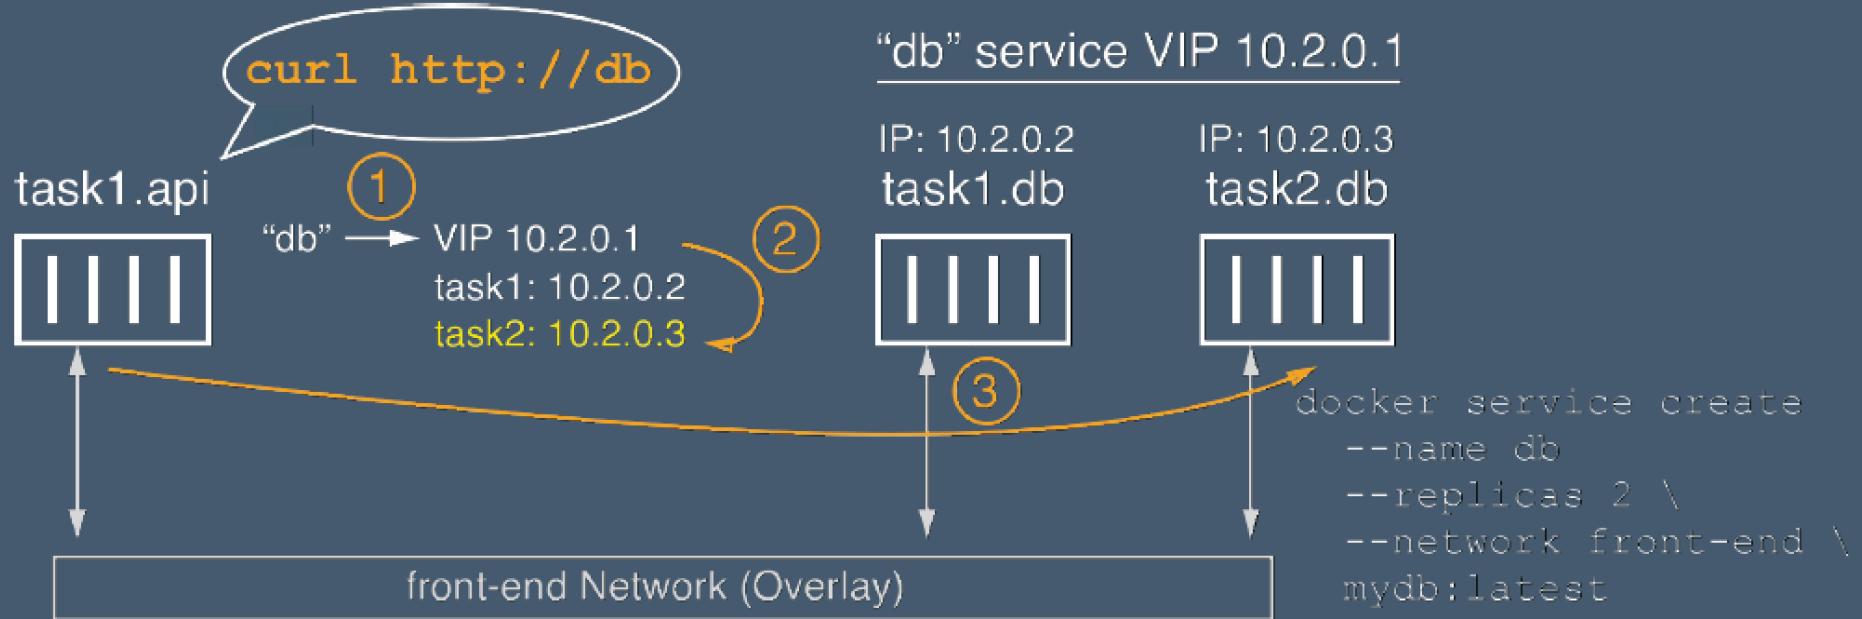
Three questions:

- Is the request originator **internal** or **external** to your cluster?
- Are the destination containers **stateless** or **stateful**?
- Are you using **Swarm** or **Kubernetes**?

ee2.1-dceu-v1.2 © 2018 Docker, Inc.

# INTERNAL / STATELESS / SWARM

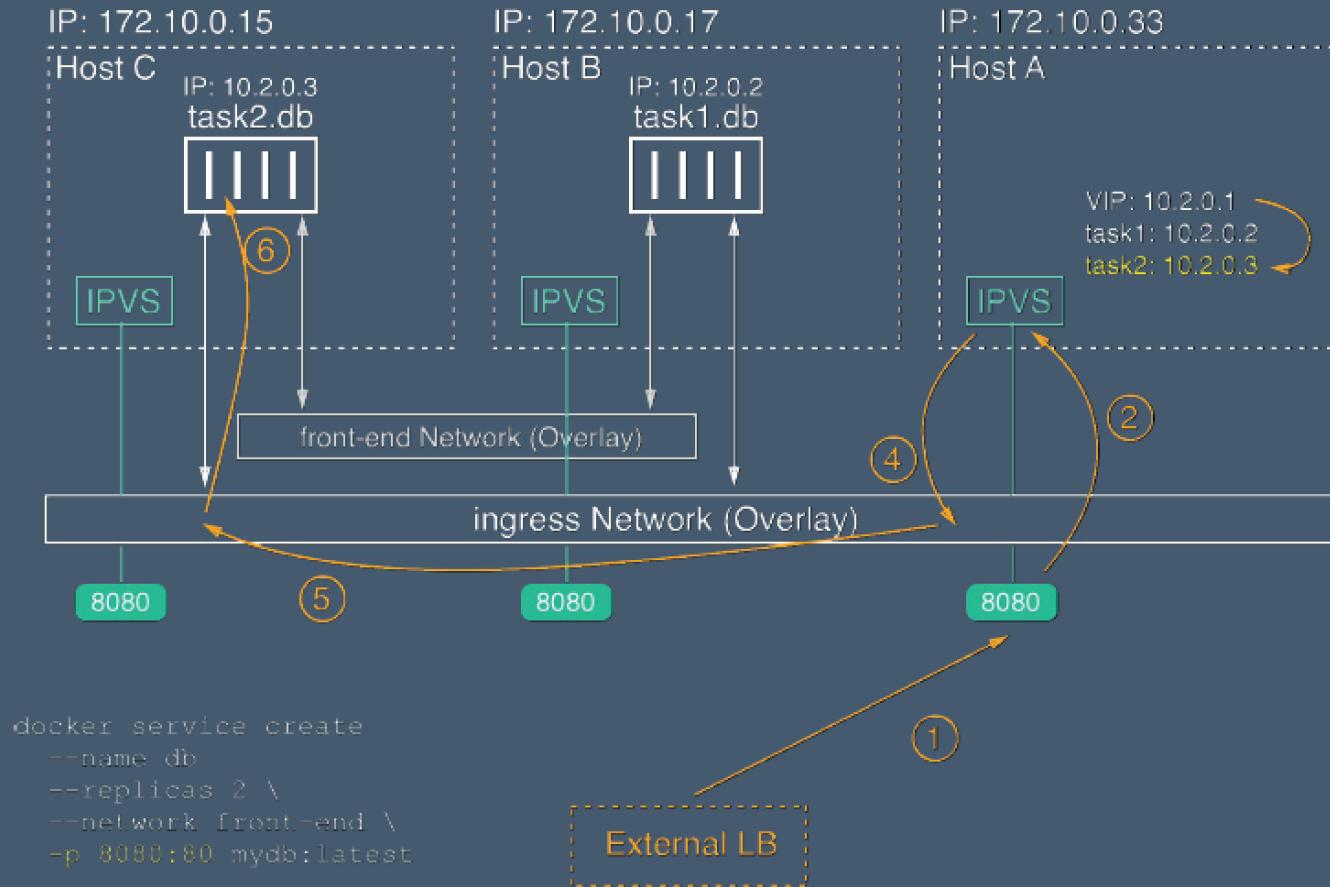
Solution: Swarm VIPs



ee2.1-dceu-v1.2 © 2018 Docker, Inc.

# EXTERNAL / STATELESS / SWARM

Solution: Swarm L4 Routing Mesh



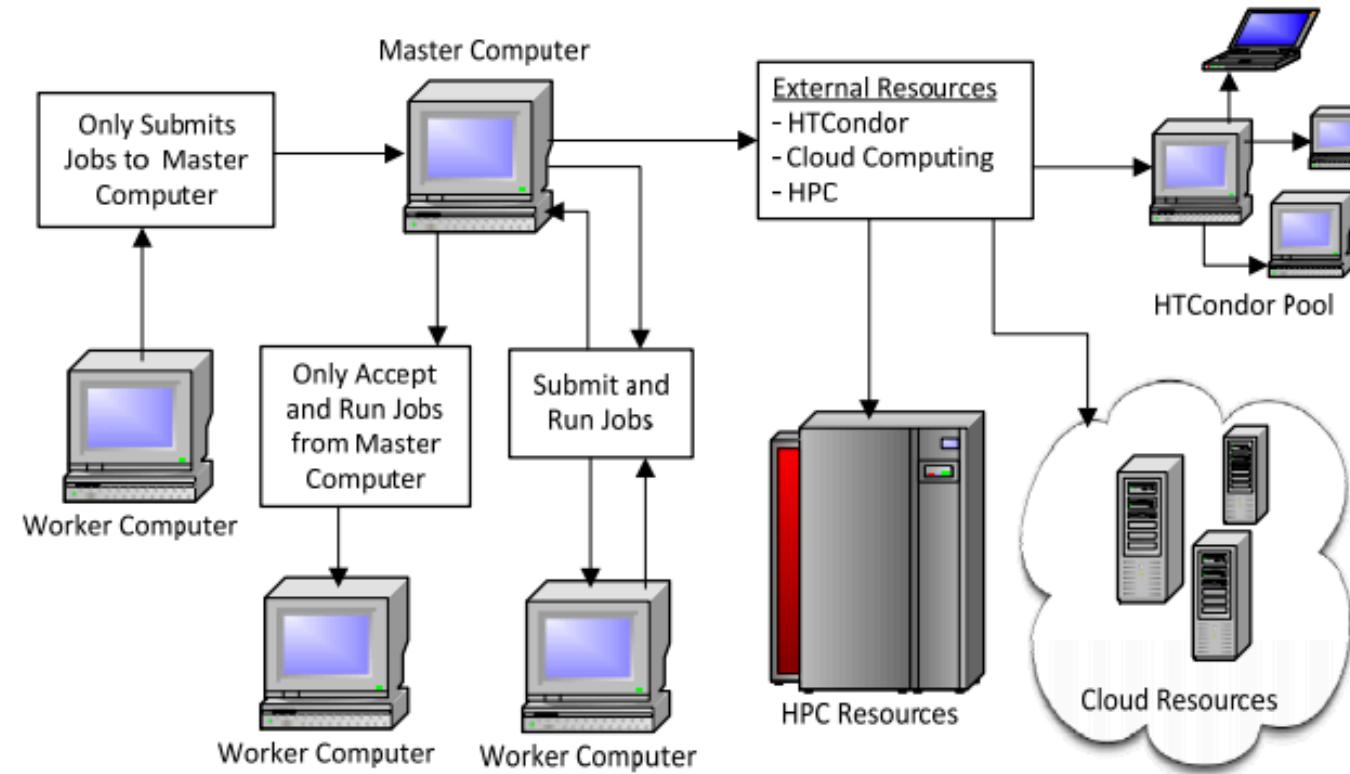
```
docker service create
--name db
--replicas 2 \
--network front-end \
-p 8080:80 mydb:latest
```

ee2.1-dceu-v1.2 © 2018 Docker, Inc.



# Use case: Containerised HPC With HTConcor and swarm

# HTCondor





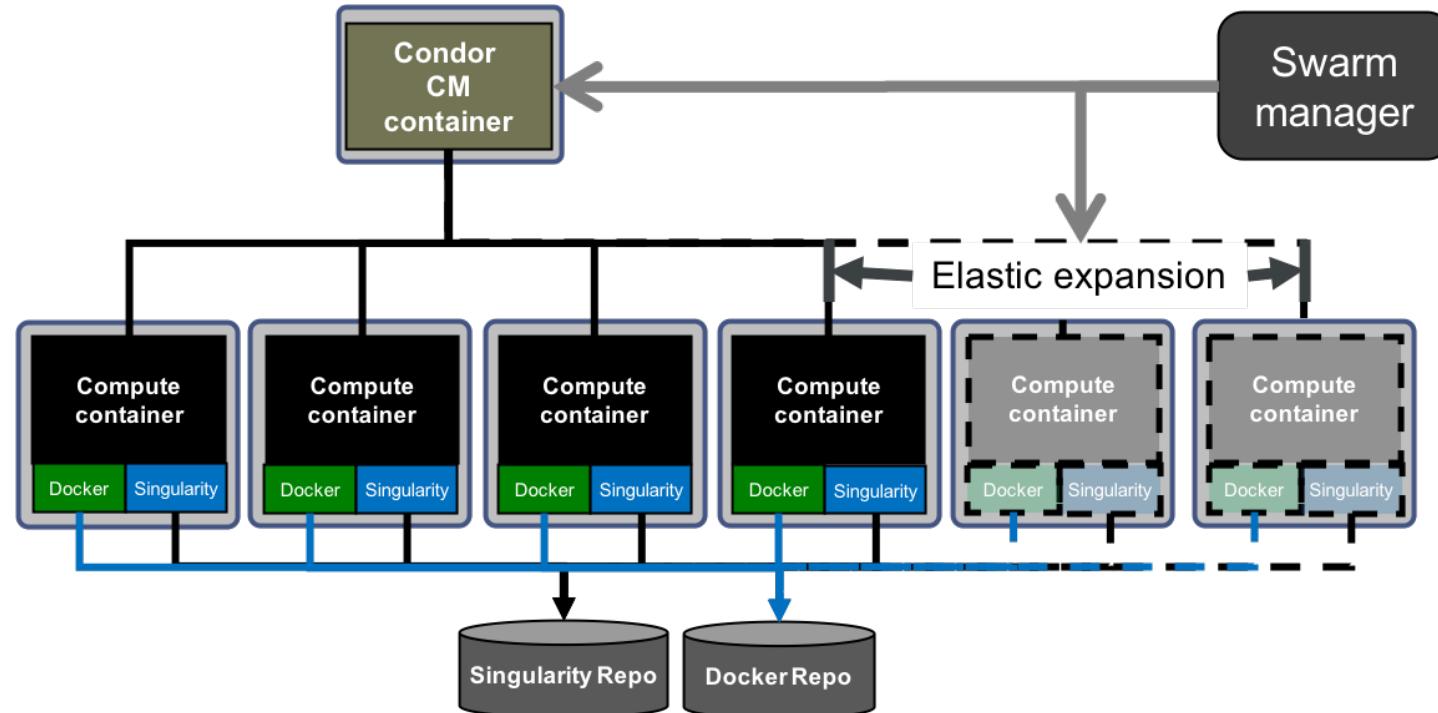
# HTCondor VM universe

```
universe = vm
executable = vmware_sample_job
log = simple.vm.log.txt
vm_type = vmware
vm_memory = 64
vmware_dir = C:\condor-test
vm_checkpoint = true
queue
```

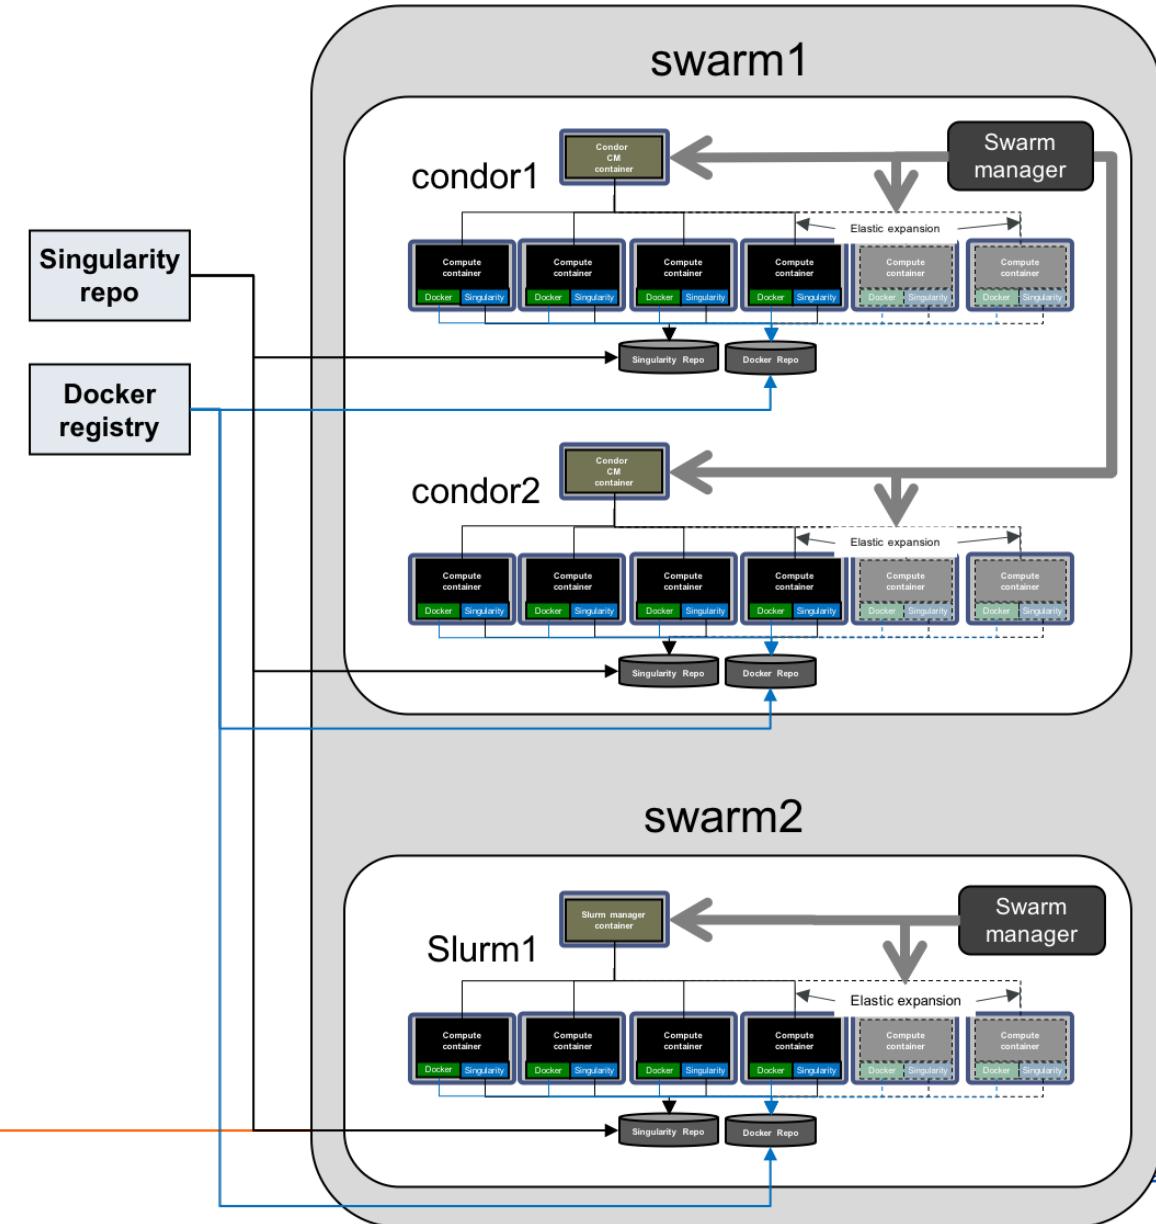
# HTCondor Docker universe

```
universe = docker
docker_image = debian
executable = /bin/cat
arguments = /etc/hosts
output = out.$(Process)
error = err.$(Process)
request_memory = 100M
queue 10
```

# Containerised HTCondor on Swarm



## Docker only cluster



ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
smzp49rwzhc6g5m3rbnwa4o41 *	swarm-1.novalocal	Ready	Active	Leader
2x9m0zr61qnm7gvnhwwjgn263	swarm-2.novalocal	Ready	Active	
n99dyxddexkr0w0wq8ysz3tw3	swarm-3.novalocal	Ready	Active	

```
[cloud-user@swarm-1 ~]$ docker service ls
```

ID	NAME	MODE	REPLICAS
	PORTS		
5rme9l2t03vi	galaxy_galaxy-htcondor	replicated	1/1
axy-htcondor:18.01b			
z6seh1nyhjep	galaxy_galaxy-htcondor-executor	replicated	2/2
axy-htcondor-executor:18.01b			
vjq5y0dd5ni5	galaxy_galaxy-htcondor-executor-big	replicated	1/1
axy-htcondor-executor:18.01b			

[cloud-user@swarm-1 ~]\$ docker service ps galaxy_galaxy-htcondor-executor   grep Running				
59zcdulus630	galaxy_galaxy-htcondor-executor.1	quay.io/bgruening/galaxy-htcondor-executor:18.01b		swarm-3
.novalocal	Running	Running	2 months ago	
4dygqif85wi2	galaxy_galaxy-htcondor-executor.2	quay.io/bgruening/galaxy-htcondor-executor:18.01b		swarm-3
.novalocal	Running	Running	2 months ago	



```
[cloud-user@swarm-1 ~]$ docker service scale galaxy_galaxy-htcondor-executor=5  
galaxy_galaxy-htcondor-executor scaled to 5  
overall progress: 5 out of 5 tasks
```

```
1/5: running [=====>]  
2/5: running [=====>]  
3/5: running [=====>]  
4/5: running [=====>]  
5/5: running [=====>]
```

```
verify: Service converged
```

```
[cloud-user@swarm-1 ~]$ docker service ps galaxy_galaxy-htcondor-executor | grep Running  
59zcdulus630      galaxy_galaxy-htcondor-executor.1      quay.io/bgruening/galaxy-htcondor-executor:18.01b      swarm-3  
.novalocal  Running          Running 2 months ago  
4dygqif85wi2      galaxy_galaxy-htcondor-executor.2      quay.io/bgruening/galaxy-htcondor-executor:18.01b      swarm-3  
.novalocal  Running          Running 2 months ago  
t1cadorzttuh     galaxy_galaxy-htcondor-executor.3      quay.io/bgruening/galaxy-htcondor-executor:18.01b      swarm-2  
.novalocal  Running          Running 15 seconds ago  
av7mt5t7i6jk     galaxy_galaxy-htcondor-executor.4      quay.io/bgruening/galaxy-htcondor-executor:18.01b      swarm-2  
.novalocal  Running          Running 15 seconds ago  
920zv43mmaox     galaxy_galaxy-htcondor-executor.5      quay.io/bgruening/galaxy-htcondor-executor:18.01b      swarm-2  
.novalocal  Running          Running 15 seconds ago
```



```
[cloud-user@swarm-3 ~]$ docker exec 53463399ae3a condor_status
```

Name	OpSys	Arch	State	Activity	LoadAv	Mem	ActvtyTime
2b928ae59394	LINUX	X86_64	Unclaimed	Idle	0.160	1024	0+00:00:03
5ed8939b3952	LINUX	X86_64	Unclaimed	Idle	0.000	2048	77+05:43:00
2073aae2920d	LINUX	X86_64	Unclaimed	Idle	0.160	1024	0+00:00:03
a4740170e2cf	LINUX	X86_64	Unclaimed	Idle	0.000	1024	77+05:46:51
c3d66659047a	LINUX	X86_64	Unclaimed	Idle	0.000	1024	77+05:46:55
c55e3685352c	LINUX	X86_64	Unclaimed	Idle	0.160	1024	0+00:00:03
Total Owner Claimed Unclaimed Matched Preempting Backfill Drain							
X86_64/LINUX	6	0	0	6	0	0	0
Total	6	0	0	6	0	0	0





# Kubernetes and HPC



Daniel Gruber  
The UberCloud

# Kubernetes Intro: Native Batch Job Submission

```
kubectl apply -f job.yaml
```

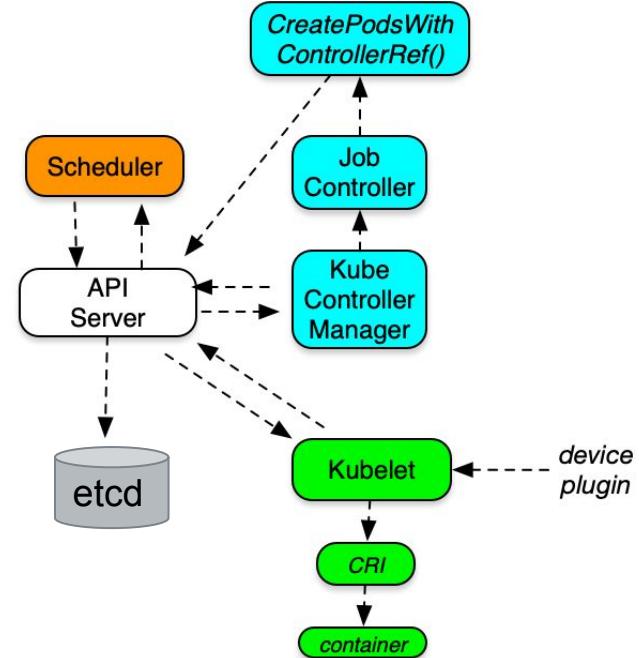
```
1  apiVersion: batch/v1
2  kind: Job
3  metadata:
4    name: sleep
5  spec:
6    template:
7      spec:
8        containers:
9          - name: sleep
10            image: busybox
11            command: ["sleep", "100"]
12            restartPolicy: Never
```



# Kubernetes Intro: Native Batch Job Submission

kubectl apply -f job.yaml

```
1  apiVersion: batch/v1
2  kind: Job
3  metadata:
4    name: sleep
5  spec:
6    template:
7      spec:
8        containers:
9          - name: sleep
10         image: busybox
11         command: ["sleep", "100"]
12         restartPolicy: Never
```



Deployment of  
more complex  
applications:  
*helm,  
Operators*

# Kubernetes and HPC

HPC Requirements	Kubernetes features
➤ Fast job submission times	❖ <i>kubectl</i> , API
➤ Multi-node jobs (distributed memory)	❖ No direct MPI Support
➤ Support for GPUs, FPGAs	❖ GPUs supported via <i>device plugin</i>
➤ Easily accessible by HPC applications	❖ AI
➤ Job queueing capabilities	❖ Pods, since 1.9
➤ Job prioritization	❖ 1.14
➤ Job pre-emption	❖ 1.14

# Kubernetes and HPC (continued)

HPC Requirements	Kubernetes features
➤ Advance reservation	❖ No
➤ License management capabilities	❖ No, Scheduler Extenders
➤ Support for Infiniband, Omnipath	❖ device plugin
➤ Full hardware utilization	❖ More daemons, typically virtualized, CPU manager
➤ Resource control	❖ Memory and CPU resources, Resource quotas
➤ Accounting and reporting	❖ Yes
➤ Storage	❖ PVs and PVCs, storageclass, hostpath

# References

- Kubeflow: <https://github.com/kubeflow/kubeflow>
- MPI Operator: <https://github.com/kubeflow/mpi-operator>
- Poseidon: <https://github.com/kubernetes-sigs/poseidon>
- CPU Manager: <https://kubernetes.io/docs/tasks/administer-cluster/cpu-management-policies/>
- Kube batch: <https://github.com/kubernetes-sigs/kube-batch>
- Volcano: <https://github.com/volcano-sh/volcano>
- Slurm Operator: <https://github.com/sylabs/slurm-operator>
- Lustre: <https://github.com/kvaps/kube-lustre>
- DRMAA2: <https://github.com/dgruber/drmaa2os>
- qsub for k8s: <https://github.com/dgruber/qsub>
- RDMA: <https://github.com/Mellanox/k8s-rdma-sriov-dev-plugin>
- Univa's Command: <http://www.univa.com/products/navops.php>
- UberCloud: <https://www.theubercloud.com/>



# Thank you

Daniel.Gruber@TheUberCloud.com

Some of my strange thoughts: [www.drmaa2.org](http://www.drmaa2.org) (→ <https://www.gridengine.eu>)



<https://github.com/dgruber>



## New on the Open Source blog: Using the FSx for Lustre CSI Driver with Amazon EKS

by Sean White | on 16 APR 2019 | in [Amazon Elastic Container Service](#), [Amazon Elastic Container Service For Kubernetes](#), [Amazon FSx For Lustre](#), [Storage](#) | [Permalink](#) | [Comments](#) | [Share](#)

From time to time the AWS Storage Blog will point you to storage related blog posts from other AWS blog channels that we think you will find interesting or helpful for your storage use cases. We wanted to share a post from the [Open Source Blog](#) that highlights a walk-through of using Amazon FSx for Lustre Container Storage Interface (CSI) Driver with Amazon EKS.

The [Amazon Elastic Container Service for Kubernetes \(Amazon EKS\)](#) team has been busy building CSI drivers for all our storage solutions, including [Amazon FSx for Lustre](#): a fully-managed file system which has integrations with S3 and is optimized for compute-intensive workloads such as high-performance computing (HPC) and machine learning. Because the AWS FSx CSI Driver is potentially useful to any Kubernetes user, we've donated it to Kubernetes SIG-AWS. This blog post focuses on deploying the AWS FSx CSI driver to an Amazon EKS cluster.

You can now dynamically provision [Amazon FSx for Lustre](#) filesystems for your high performance computing workloads, and have a container automatically connected into this filesystem. FSx for Lustre provides a high-performance file system optimized for fast processing of workloads such as machine learning, high performance computing (HPC), video processing, financial modeling, and electronic design automation (EDA). These workloads commonly require data to be presented via a fast and scalable file system interface, and typically have data sets stored on long-term data stores like Amazon S3.

You can also use dynamic provisioning to consume the same persistent volume claim from multiple pods from different nodes. If you need to start your containers with a dataset automatically available, for example loading a machine learning training data set, you can even connect your StorageClass to an existing Amazon S3 bucket by using Dynamic Provisioning with Data Repository. Other examples can be found in the FSx for Lustre CSI Driver documentation.

To learn more, see the full [AWS Open Source Blog post](#).

<https://aws.amazon.com/blogsopensource/using-fsx-lustre-csi-driver-amazon-eks/>



# <https://github.com/kubernetes-sigs/aws-fsx-csi-driver>

CSI Driver of AWS FSx for Lustre <https://aws.amazon.com/fsx/lustre/>

aws fsx csi kubernetes k8s-sig-aws

91 commits 1 branch 2 releases 11 contributors Apache-2.0

Branch: master New pull request Find File Clone or download

k8s-ci-robot Merge pull request #64 from benoitbayo/patch-1 ... Latest commit d20c305 on May 1

File	Description	Time Ago
.github	Update issue templates (#31)	3 months ago
cmd	Fix klog init flag	4 months ago
deploy/kubernetes	adding single manifest for driver	3 months ago
docs	Add iam policy for FSx driver	2 months ago
examples/kubernetes	updated specs and README for fsx size	2 months ago
hack	Remove install lustre client hack script	3 months ago
pkg	Fix typo	2 months ago
tests/sanity	Implement dynamic provisioning for FSx for Lustre PV (#14)	4 months ago
.gitignore	Change static and dynamic example to use ReadWriteMany by default	3 months ago
.travis.yml	`aws/aws-fsx-csi-driver` -> `kubernetes-sigs/aws-fsx-csi-driver`	3 months ago
CHANGELOG-0.x.md	`aws/csi-driver-amazon-fsx` -> `kubernetes-sigs/aws-fsx-csi-driver`	3 months ago
CONTRIBUTING.md	Add Kubernetes code of conduct OWNER file	3 months ago
Dockerfile	Update dockerfile to install official lustre client (#36)	3 months ago
LICENSE	Creating initial file from template	6 months ago
Makefile	Update manifest to use official amazon repo for container image (#37)	3 months ago
OWNERS	Add OWNERS, SECURITY_CONTACTS and CoC files	3 months ago
SECURITY_CONTACTS	Add OWNERS, SECURITY_CONTACTS and CoC files	3 months ago
THIRD-PARTY	Support s3 data repository in dynamic provision (#33)	3 months ago
code-of-conduct.md	Add OWNERS, SECURITY_CONTACTS and CoC files	3 months ago
go.mod	Support s3 data repository in dynamic provision (#33)	3 months ago
go.sum	Support s3 data repository in dynamic provision (#33)	3 months ago

README.md

build passing coverage 48%

**WARNING:** This driver is currently an ALPHA release. This means that there may potentially be backwards compatible breaking changes moving forward. Do NOT use this driver in a production environment in its current state.

**DISCLAIMER:** This is not an officially supported Amazon product





# Using K8s Operators for RDMA Workloads

Dror Goldenberg, Liel Shoshan - Mellanox Technologies

ISC Container Workshop Frankfurt, June 2019

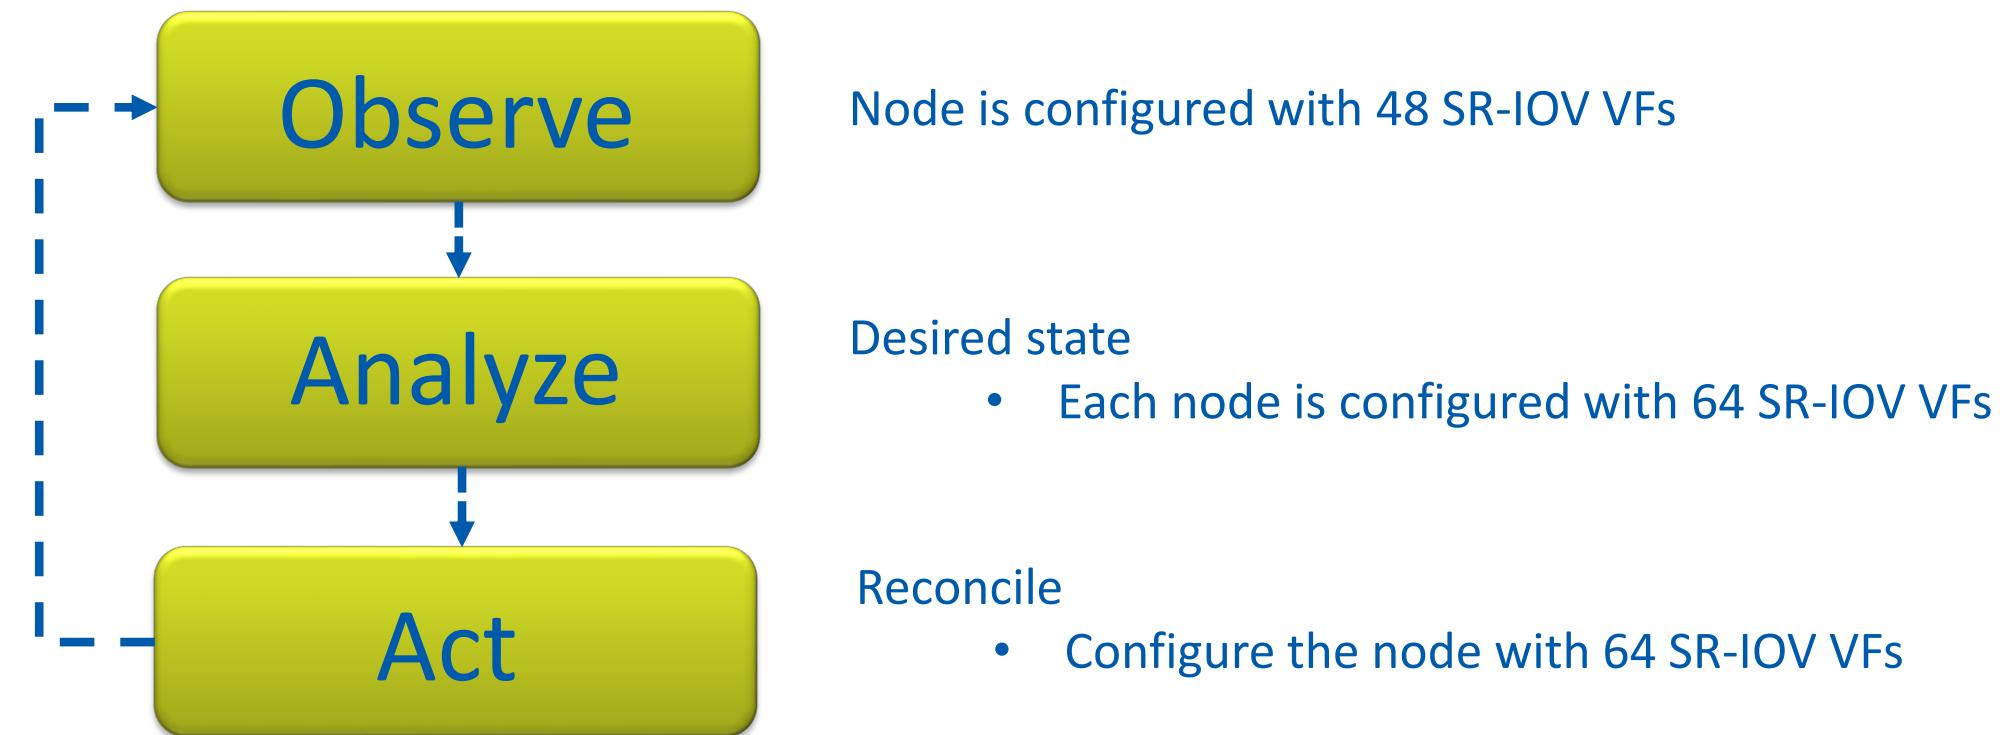


# Operators



# What are Operators?

- **Operator** - method of packaging, deploying and managing a Kubernetes application
- Enables developers to
  - Extend and add new functionalities
  - Automate administration tasks as if they were a native Kubernetes component
- Operator = set of application-specific custom controllers



# Operators

## ■ Why are Operators better than other tools out there?

- Using standard Kubernetes tools, CLI and API
- Can monitor the cluster, change pods/services, scale up/down and call endpoints of the running applications
- Smarter and more tailored than generic tools

## ■ Who builds an Operator?

- Best built by those that are experts in the “business logic” of installing, running and upgrading an application
- Creation of an Operator typically starts by automating an application installation and self-service provisioning capabilities
- Evolves to take on more complex automation

# K8s RDMA Solution



# RDMA/RoCE in K8s

- SR-IOV based solution
  - Each Container is assigned with an SR-IOV VF
  - Consist of the following K8s plugins
    - [SR-IOV Device Plugin for Kubernetes on GitHub](#)
      - SR-IOV VF provisioning
    - [SR-IOV CNI Plugin for Kubernetes on GitHub](#)
- Virtual networking in shared device mode
  - An IB device is shared between different pods



## kubernetes

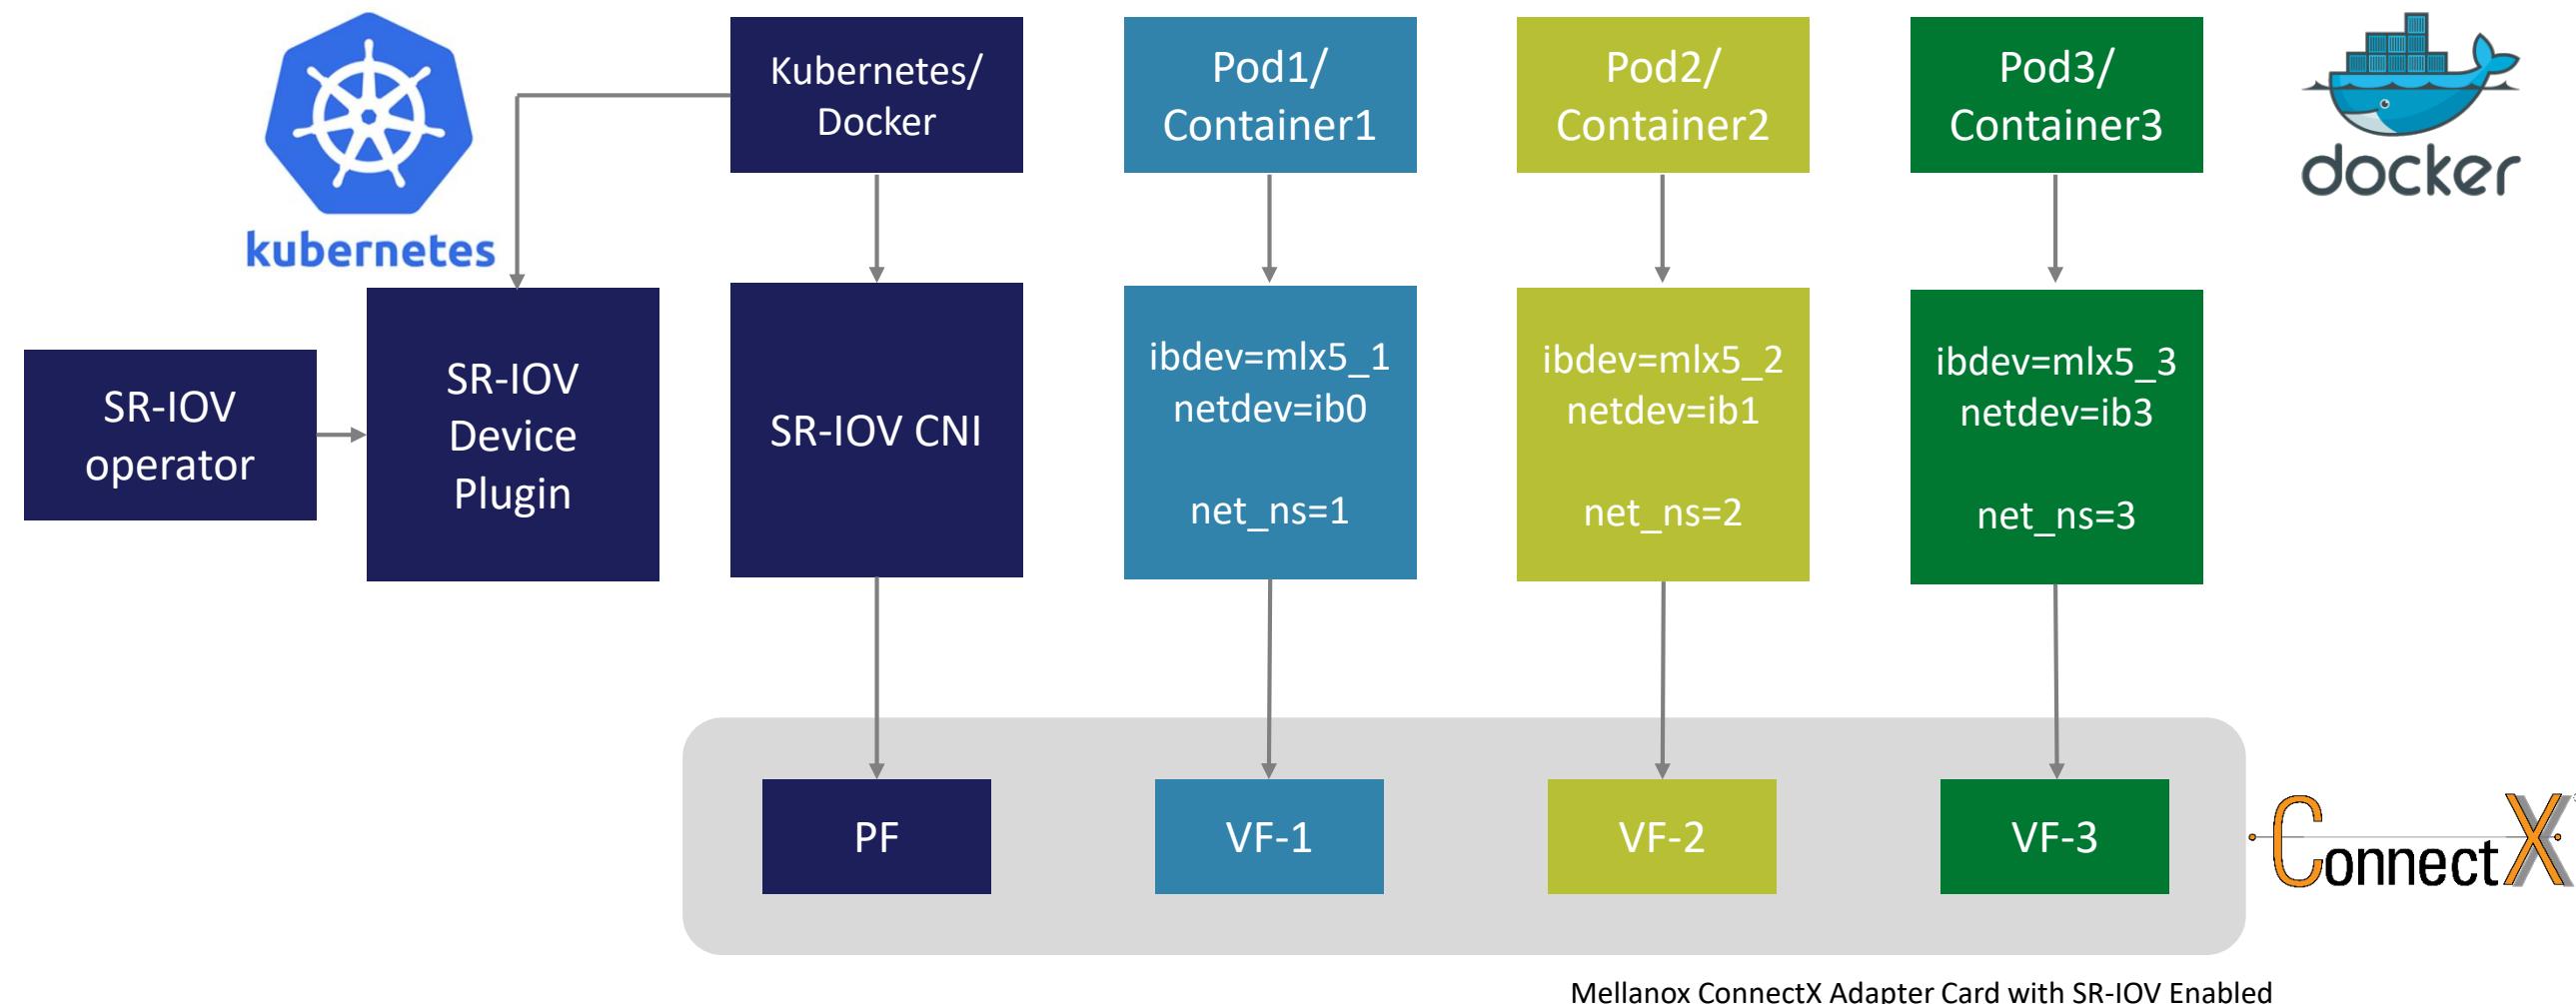


**INFINIBAND™**  
TRADE ASSOCIATION



# RDMA at the Container Level using SR-IOV

## Technical Overview

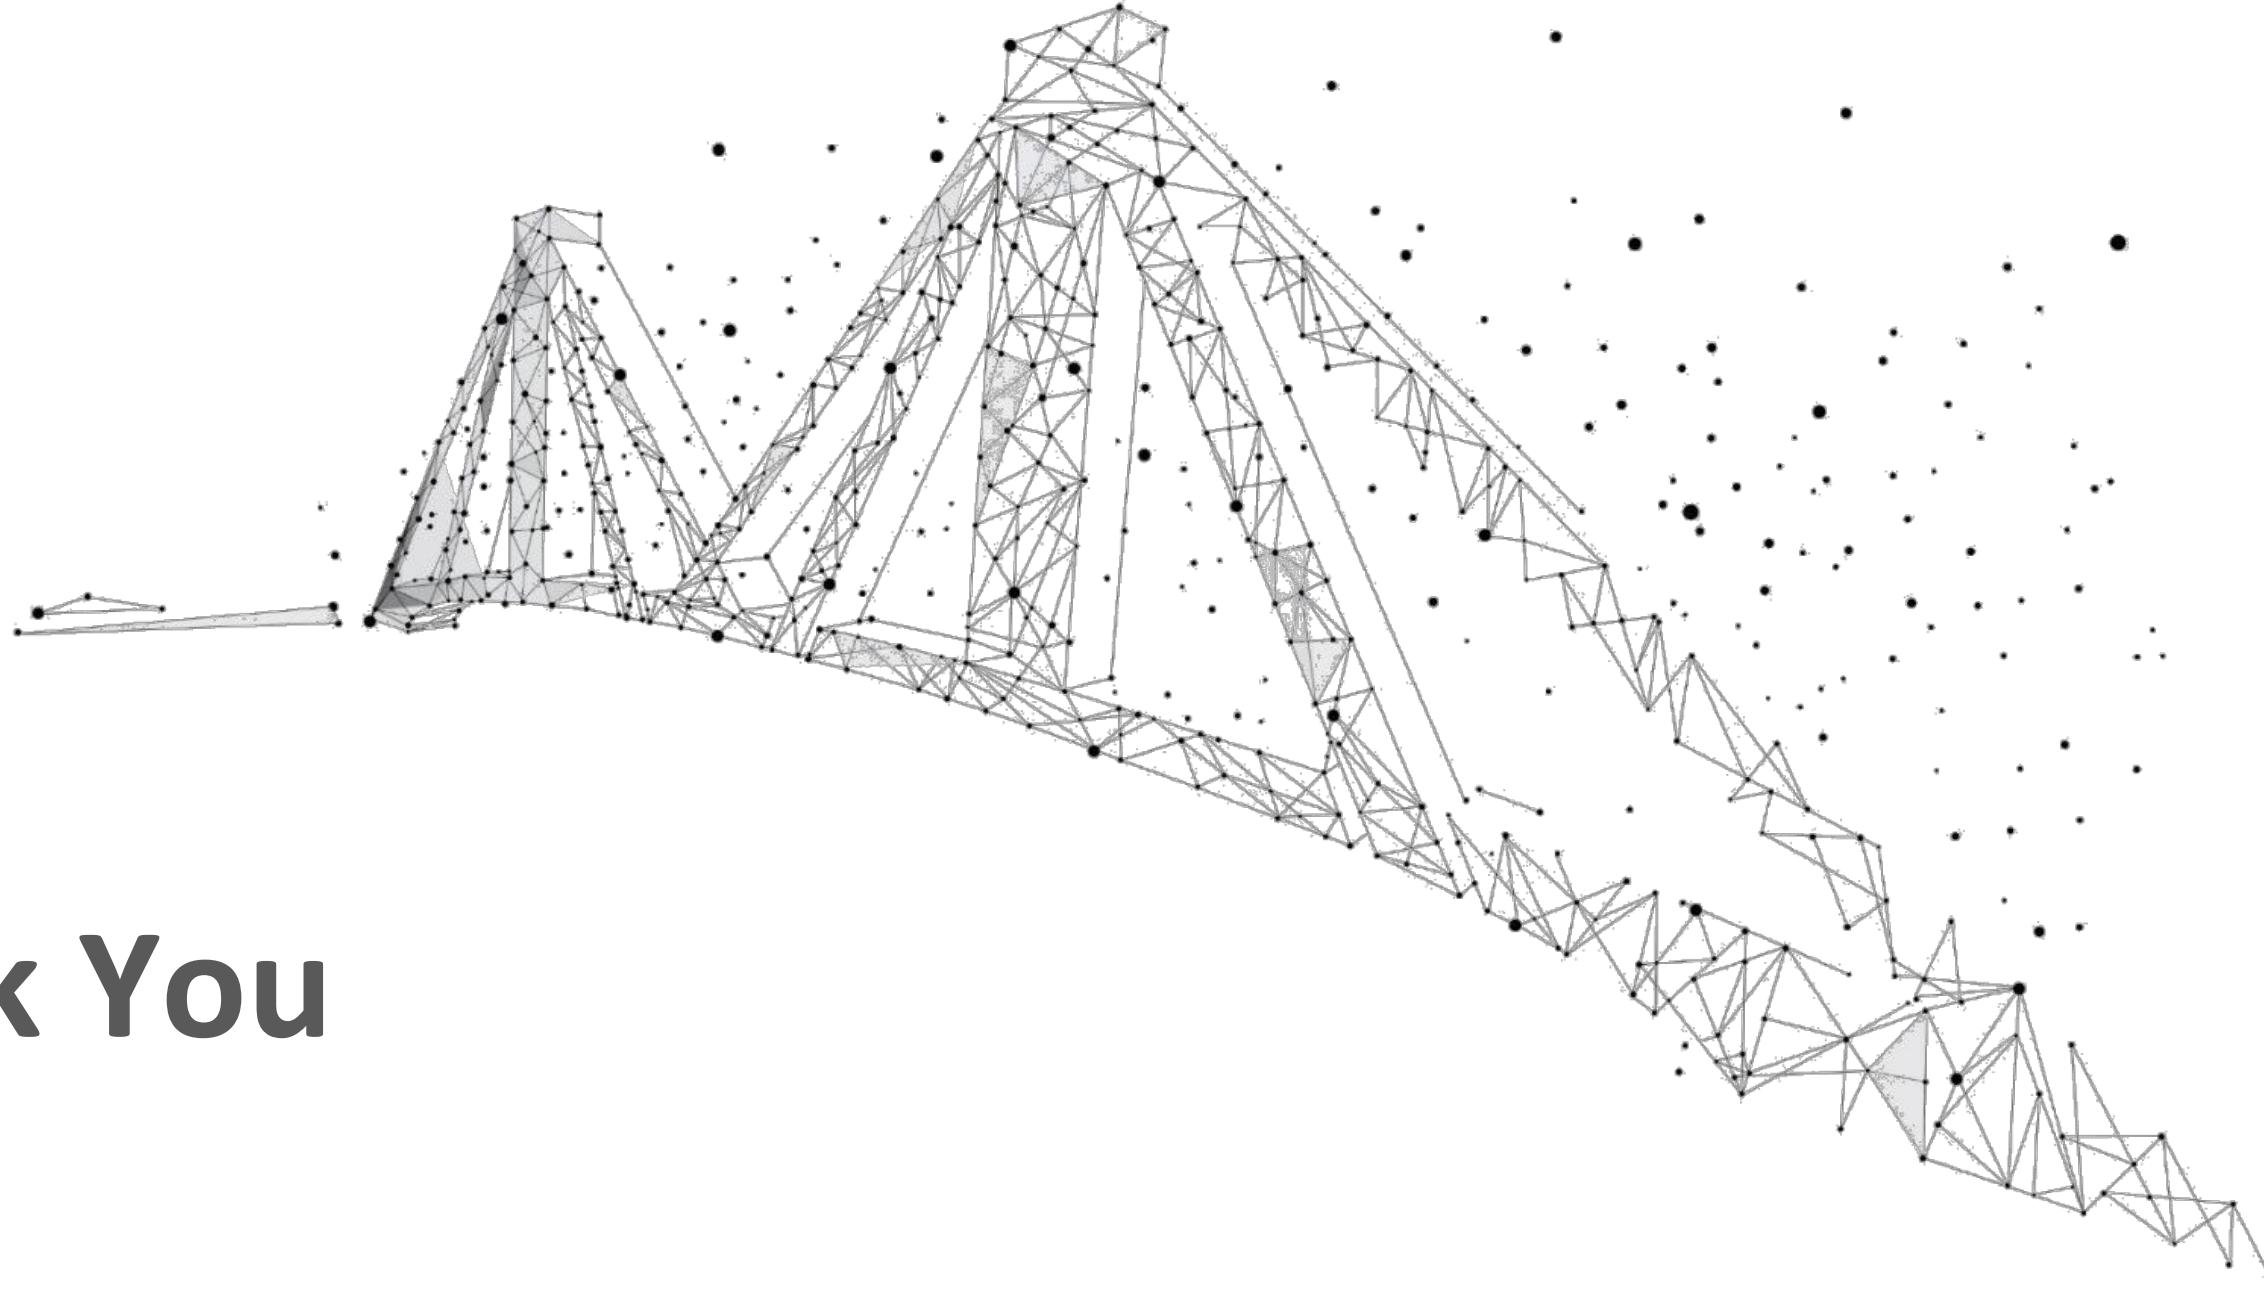


- Every container/POD has an IB device (mlx5\_1,2,3) and netdevice

# Operators for RDMA Workloads

- Why do we need operators for running RDMA workloads?
  - Running RDMA/RoCE workloads requires several configurations
  - Automating these configuration allows ease of use
- Required configurations
  - Driver installation
  - Link state configuration – Infiniband/Ethernet
- For SR-IOV based solution:
  - FW configuration
    - SR-IOV enable
    - Setting Number of VF's
  - VFs activation
  - GUID and MAC address configuration





A large, abstract geometric structure composed of numerous small triangles and dots, forming a complex, winding shape that resembles a mountain range or a network. The structure is primarily grey with some darker dots and shaded areas. It is positioned in the upper half of the slide, with its base extending towards the bottom right corner.

**Thank You**





## **Singularity, SLURM, and Kubernetes - 5 min**

---

20 June 2019

Michael Bauer - HPC Container Workshop ISC19

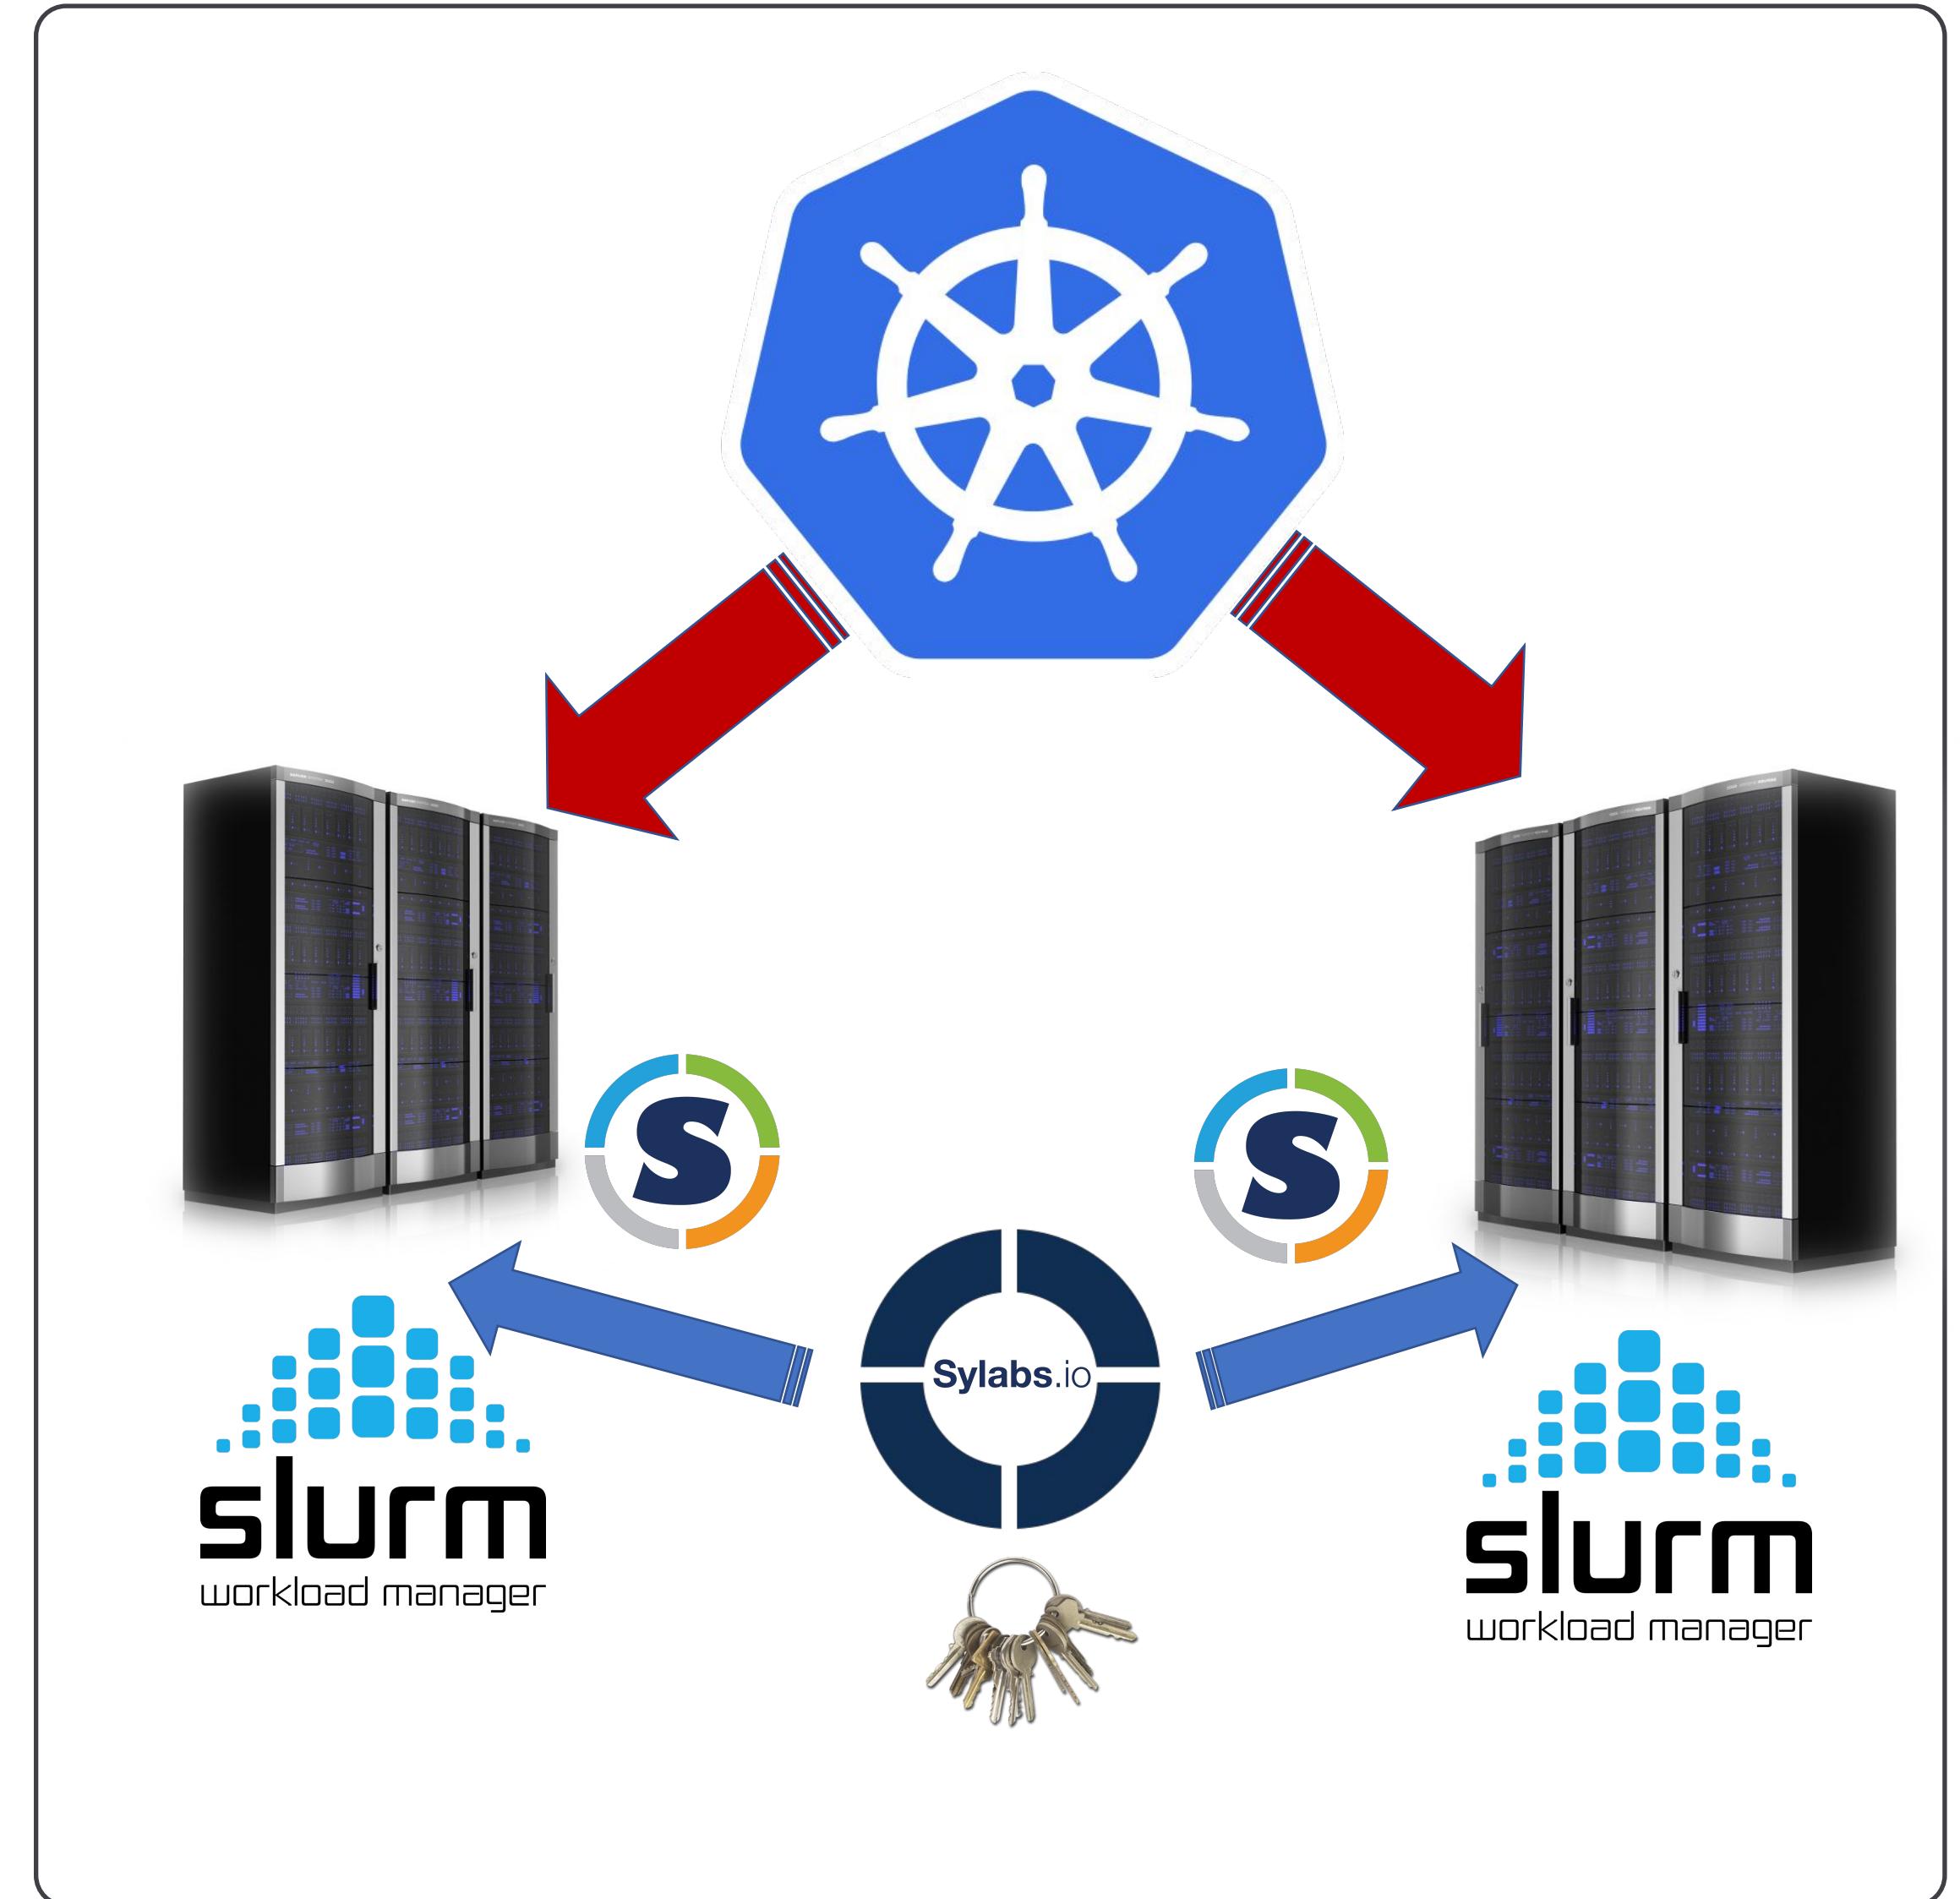
# Running Singularity with K8s

- Singularity-CRI now in **v1.0.0-alpha.3** release
  - *Please try and provide feedback!*
- Run all workloads (K8s microservices, HPC, EPC, etc...) with one unifying runtime
- Native support for **GPUs, Infiniband/RDMA** (*coming soon*), other HPC-oriented hardware via **Container Device Plugin**

# Multi-Cluster Scheduling

Kubernetes - scheduling to multiple HPC clusters - using Singularity containers to distribute the workloads.

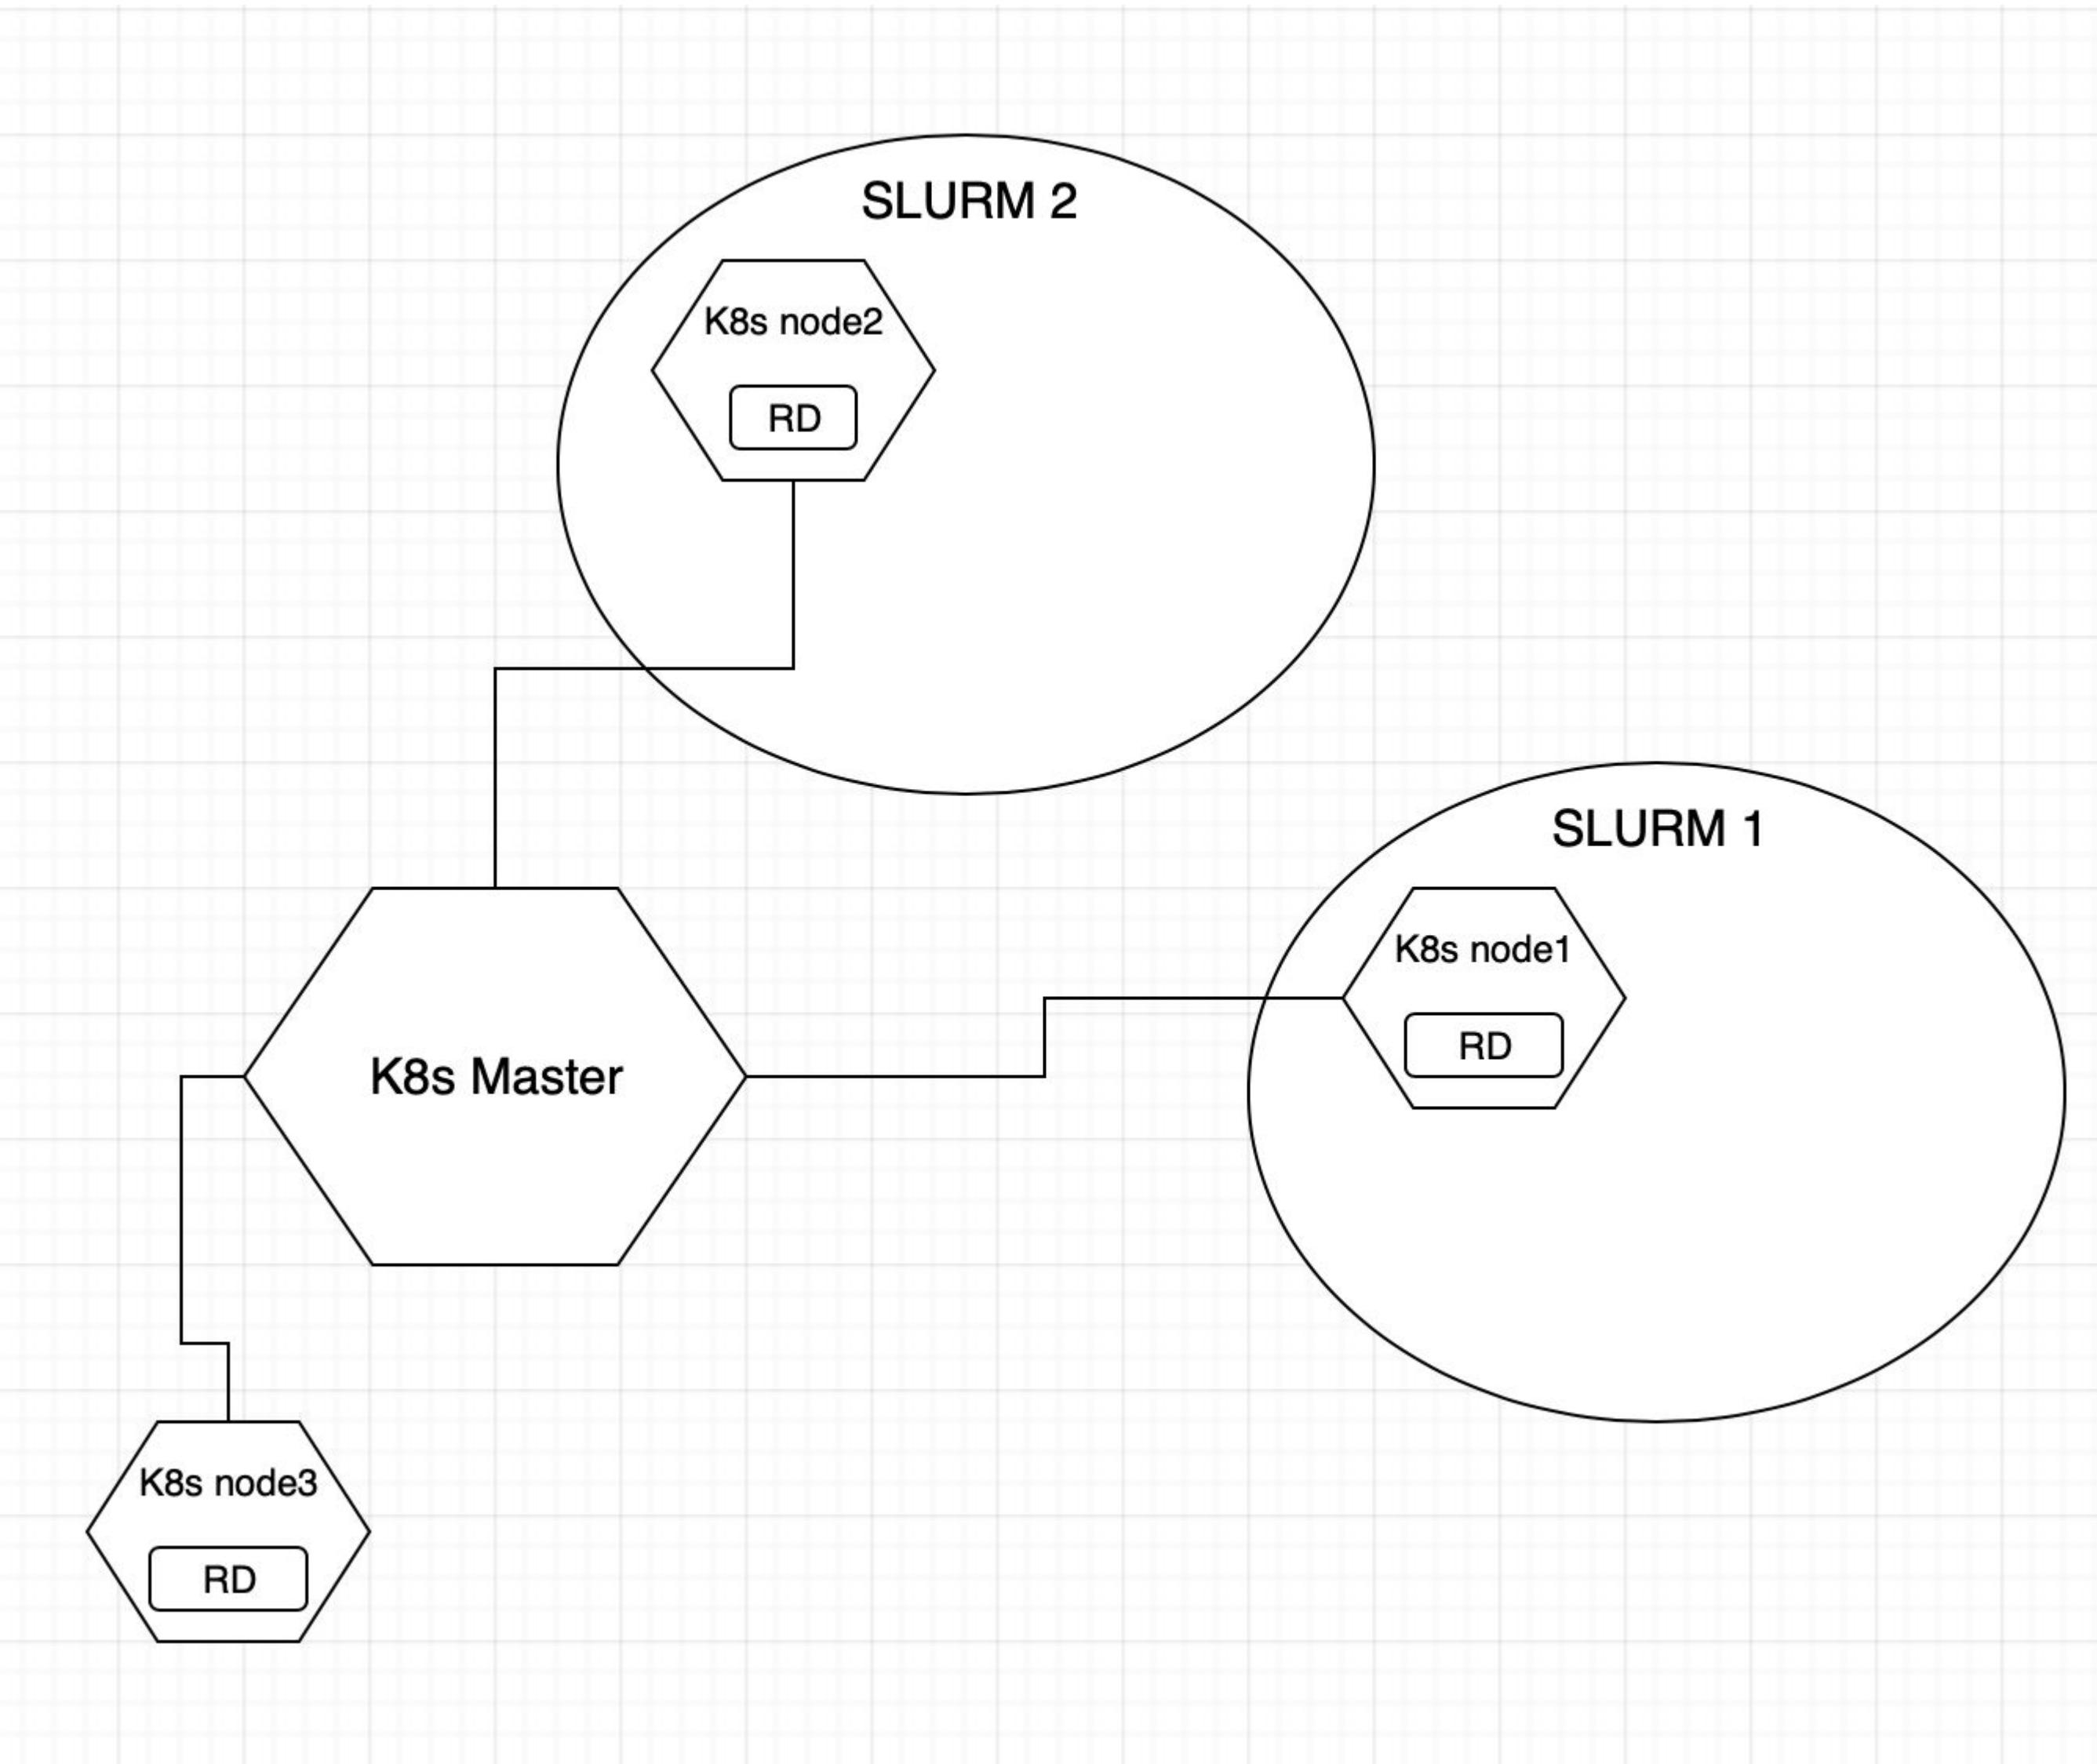
All as unprivileged user.



# Multi-Cluster Scheduling to SLURM via K8s

- Submit SLURM job request to Kubernetes server using new job kind “SlurmJob”
- Custom K8s Resource Daemon lives on each SLURM master node, lets K8s understand what resources exist
- Use default or custom K8s scheduling algorithm to dispatch jobs to different clusters

# Singularity, SLURM, K8s



```
apiVersion: slurm.sylabs.io/v1alpha1
kind: SlurmJob
metadata:
  name: cow
spec:
  batch: |
    #!/bin/sh
    ##SBATCH --nodes=1 --cpus-per-task=1
    srun singularity pull library://sylab sed/examples/lolcow
    srun singularity run lolcow_latest.sif
    srun rm lolcow_latest.sif
  nodeSelector:
    containers: singularity
```

# PORTABLE CONTAINERS ORCHESTRATION AT SCALE WITH NEXTFLOW

Paolo Di Tommaso, Seqera Labs  
ISC-HPC 2019 - Frankfurt



# ENABLING TECHNOLOGY

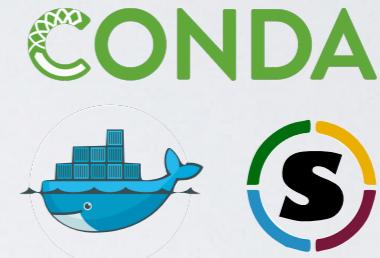
**nextflow**



code



orchestration



dependencies



deployment



sharing & reproducibility

# WHAT DO YOU MEAN?

```
# satellite sequences reported by RepeatMasker.  
zcat rmsk.txt.gz \  
| grep Satellite \  
| cut -f6,7,8 \  
| sed s,^chr,, \  
| perl -pe 's/^[\^s_]+([^\s_]+)_random/$1.1/' \  
| tr "gl" "GL" \  
| sort -k1,1N -k2,2n \  
| bgzip > hs37d5.satellite.bed.gz
```

# THE NEXTFLOW WAY

```
process filtering {
    input: file 'rmsk.txt.gz' from sequences_ch
    output: file 'hs37d5.satellite.bed.gz' into results_ch

    ...
    # satellite sequences reported by RepeatMasker.
    zcat rmsk.txt.gz \
        | grep Satellite \
        | cut -f6,7,8 \
        | sed s,^chr,, \
        | perl -pe 's/^[\s]+([^\s]+)_random/$1.1/' \
        | tr "gl" "GL" \
        | sort -k1,1N -k2,2n \
        | bgzip > hs37d5.satellite.bed.gz
    ...
}

Channel.fromPath('data/rmsk.txt.gz') | filtering | publishTo { '/path' }
```

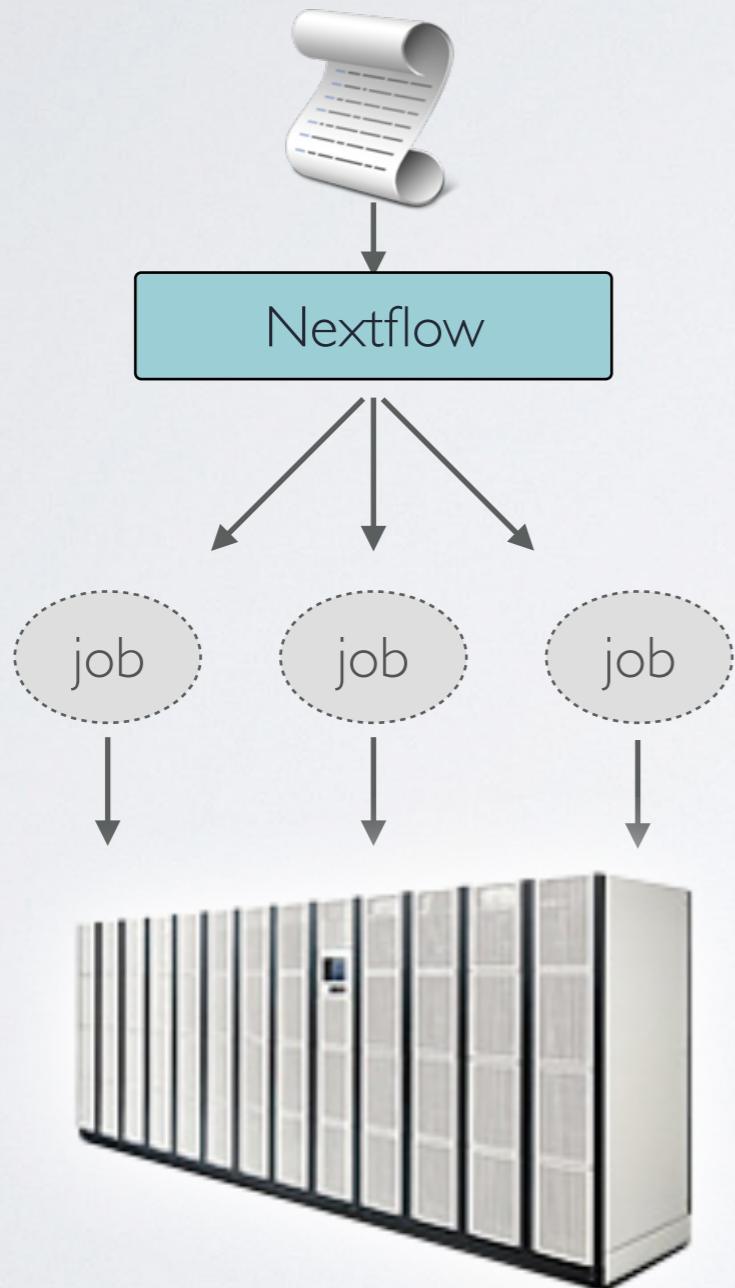
# THE NEXTFLOW WAY

```
process filtering {
    input: file 'rmsk.txt.gz' from sequences_ch
    output: file 'hs37d5.satellite.bed.gz' into results_ch

    ...
    # satellite sequences reported by RepeatMasker.
    zcat rmsk.txt.gz \
        | grep Satellite \
        | cut -f6,7,8 \
        | sed s,^chr,, \
        | perl -pe 's/^[\s]+([^\s]+)_random/$1.1/' \
        | tr "gl" "GL" \
        | sort -k1,1N -k2,2n \
        | bgzip > hs37d5.satellite.bed.gz
    ...
}

Channel.fromPath('data/*.txt.fq') | filtering | publishTo { '/path' }
```

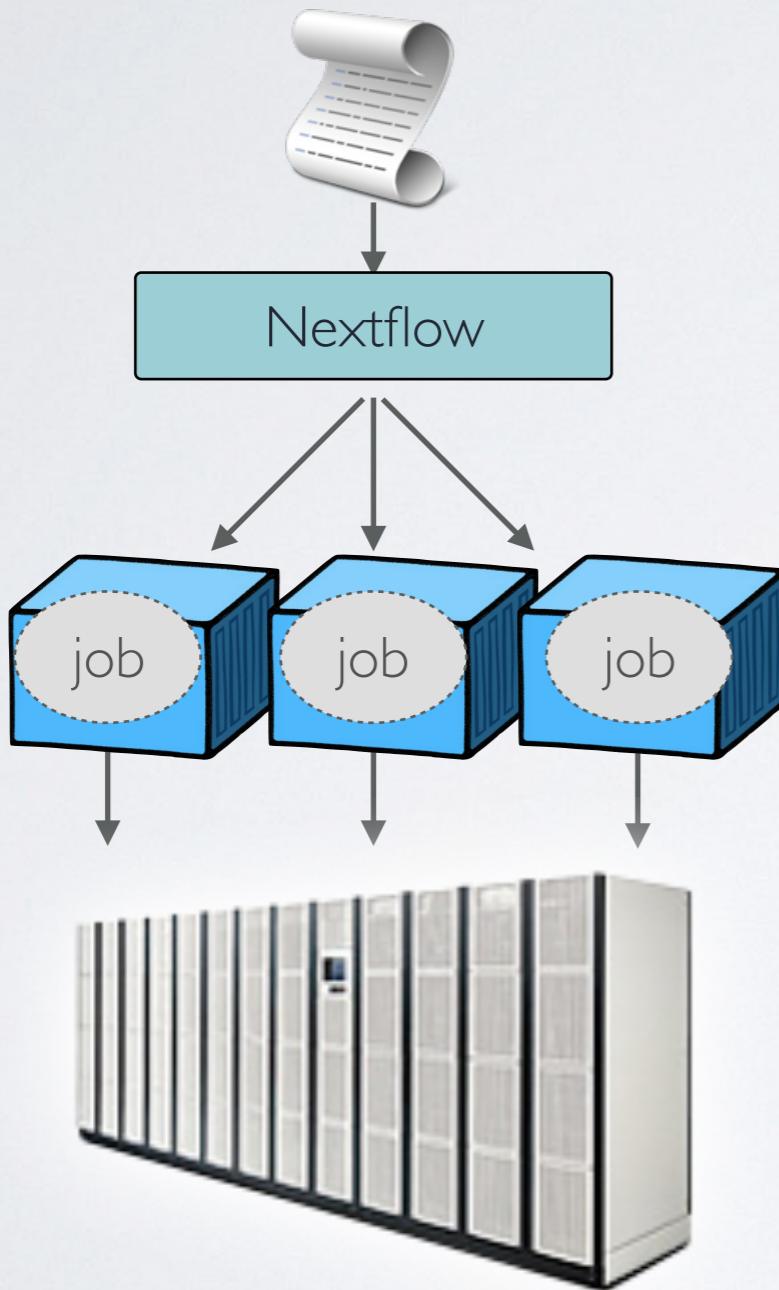
# CONTAINERISATION



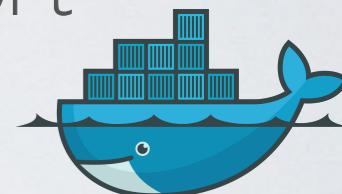
- Nextflow envisioned the use of software containers to fix computational reproducibility
- Mar 2014 (ver 0.7), support for Docker
- Dec 2016 (ver 0.23), support for Singularity



# CONTAINERISATION



- Nextflow envisioned the use of software containers to fix computational reproducibility
- Mar 2014 (ver 0.7), support for Docker
- Dec 2016 (ver 0.23), support for Singularity

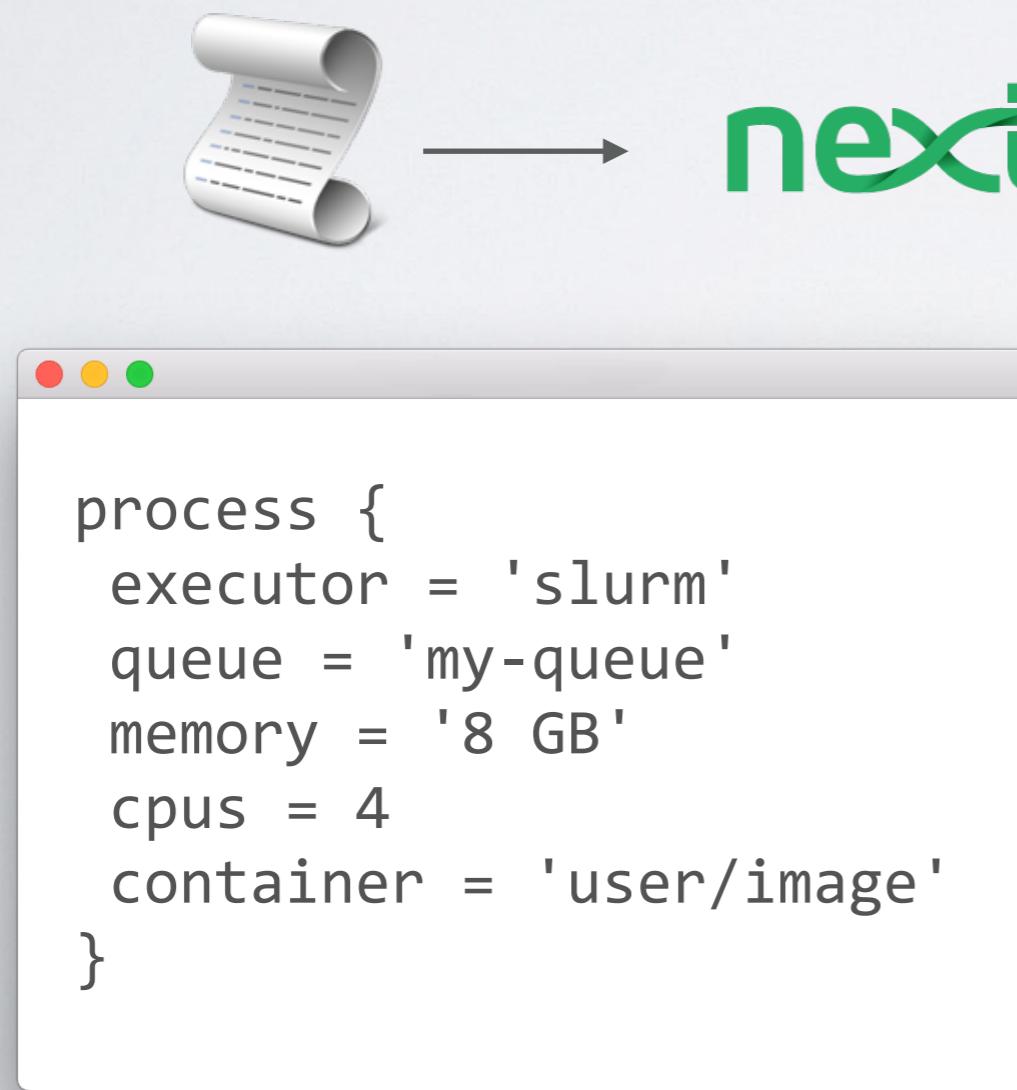


# PORTABILITY



```
nextflow run your-script.nf -with-docker your/image
```

# PORTABILITY



**nextflow**



**slurm**  
workload manager

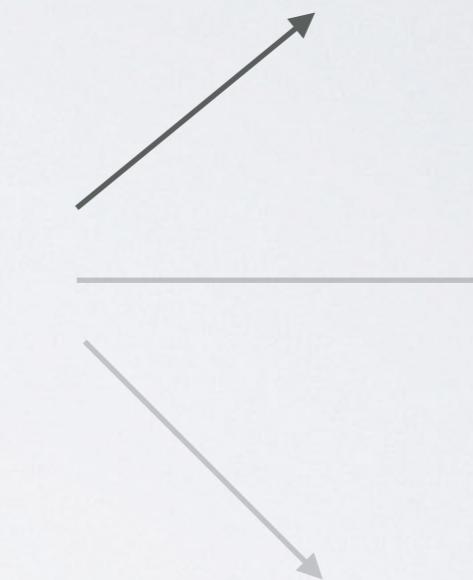
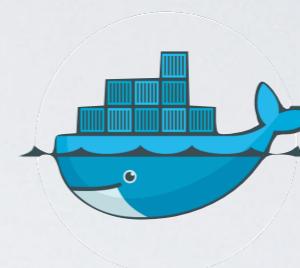


# PORTABILITY

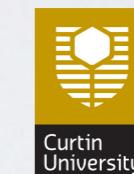
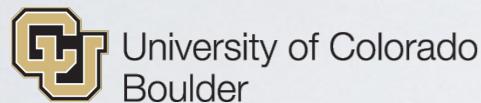


**nextflow**

```
process {  
    executor = 'awsbatch'  
    queue = 'my-queue'  
    memory = '8 GB'  
    cpus = 4  
    container = 'user/image'  
}
```



# WHO IS USING NEXTFLOW?



Weill Cornell Medical College



Agency for  
Science, Technology  
and Research  
SINGAPORE

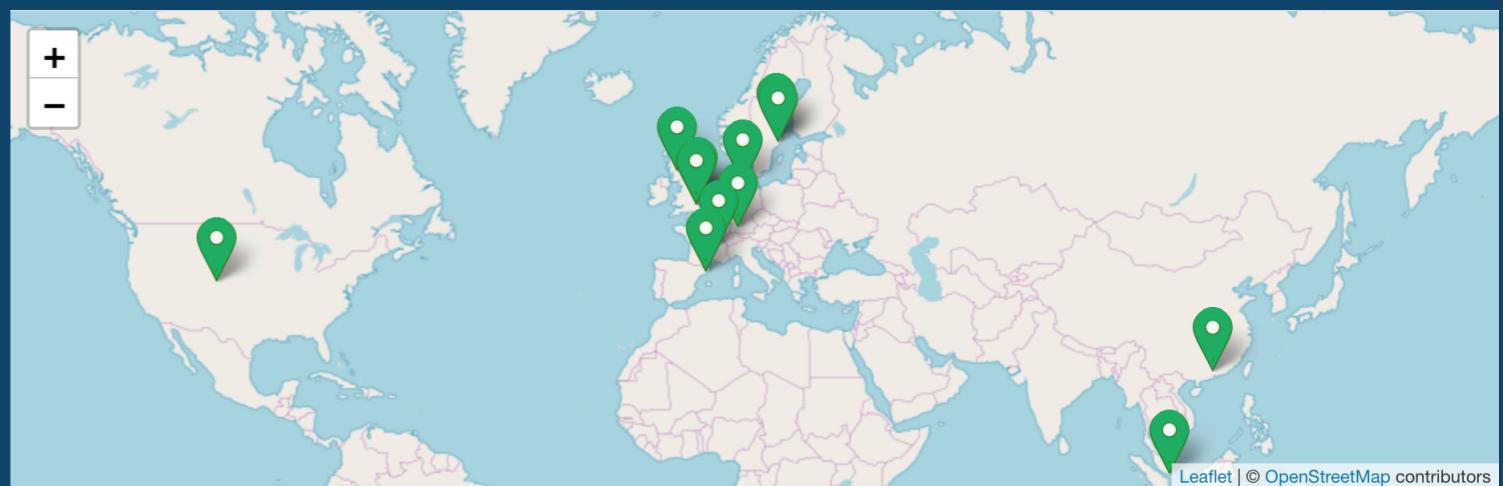
SYNTHETIC GENOMICS



# nf-core



20 pipelines    38 members    12+ institutions



CANCER  
RESEARCH  
UK

BEATSON  
INSTITUTE



NATIONAL  
GENOMICS  
INFRASTRUCTURE



Genome Institute  
of Singapore



SLU



中山大學  
腫瘤防治中心  
SUN YAT-SEN UNIVERSITY CANCER CENTER

International Agency for Research on Cancer



University  
of Colorado  
Boulder



# THANK YOU

**nextflow**

<http://nextflow.io>

 **seqera**labs

<http://seqera.io>



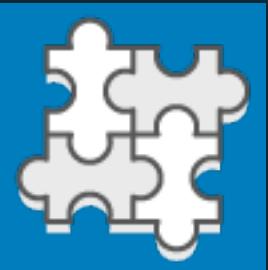
# AWS Batch

# Introducing AWS Batch



## Fully Managed

No software to install or servers to manage. AWS Batch provisions, manages, and scales your infrastructure



## Integrated with AWS

Natively integrated with the AWS Platform, AWS Batch jobs can easily and securely interact with services such as Amazon S3 and DynamoDB



## Cost-optimized Resource Provisioning

AWS Batch automatically provisions compute resources tailored to the needs of your jobs using Amazon EC2 and EC2 Spot

# AWS Batch Concepts



- **Job Queue**
- **Compute Environments**
- **Job Definitions**
- **Jobs**
  - **Single jobs vs Array jobs vs Multi-node Parallel jobs**
- **Scheduler**

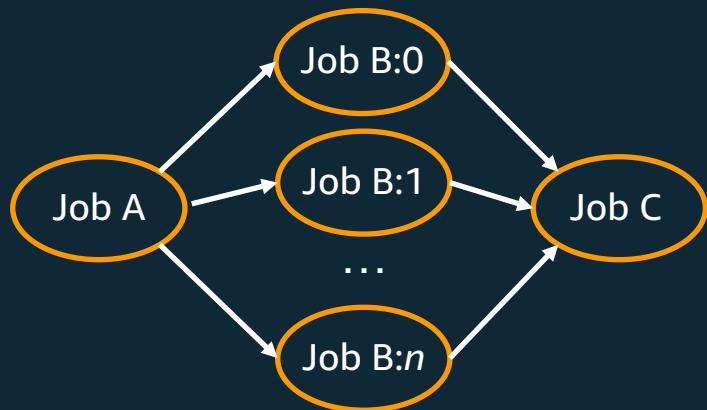
# Easily run massively parallel jobs



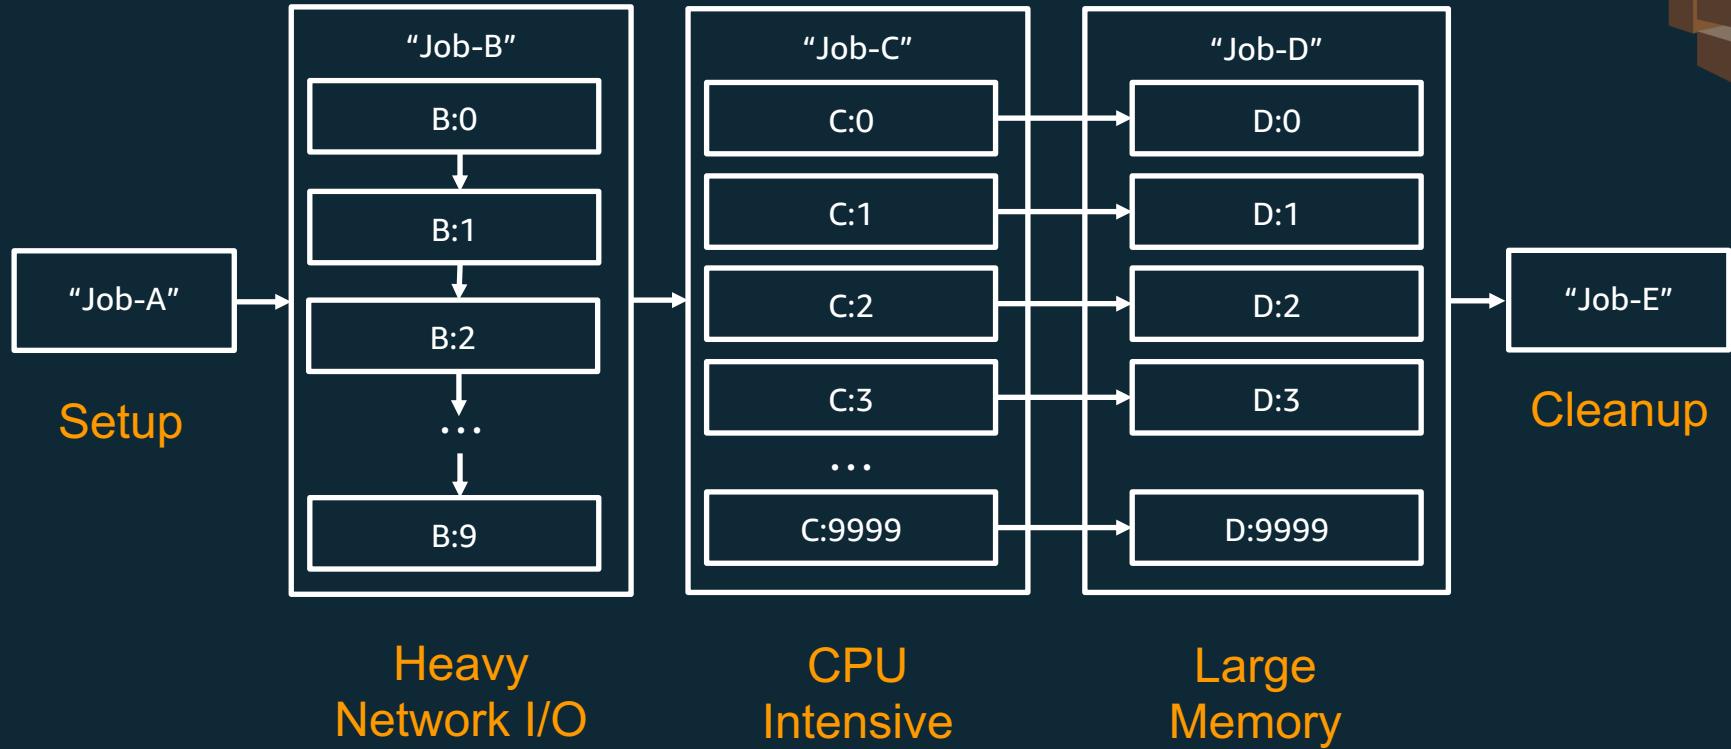
Instead of submitting a large number of independent “**simple jobs**”, we also support “**array jobs**” that run many copies of an application against an array of elements.

Array jobs are an efficient way to run:

- Parametric sweeps
- Monte Carlo simulations
- Processing a large collection of objects



# Array Job Dependency Models



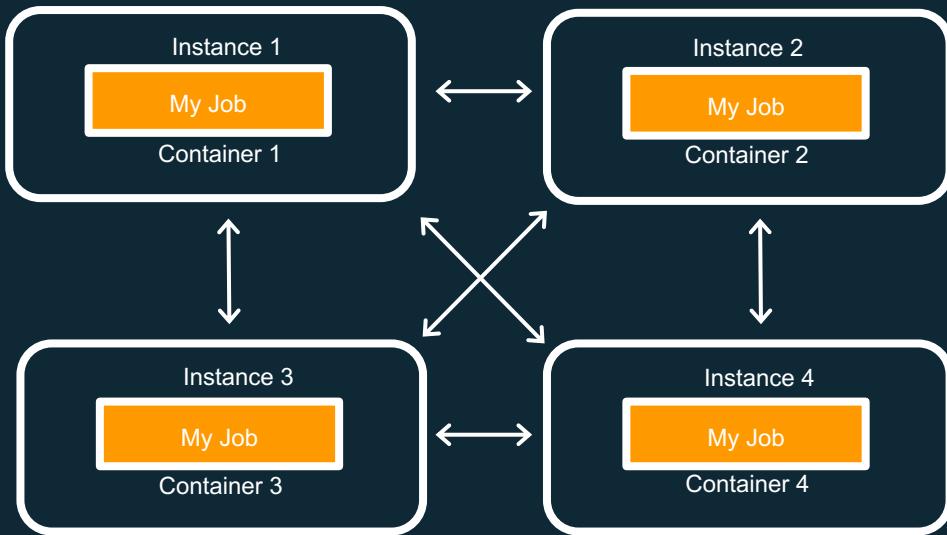
```
$ aws batch submit-job --depends-on 606b3ad1-aa31-48d8-92ec-f154bfc8215f ...
```

# Multi-node Parallel Jobs on AWS Batch

NEW

- Scale jobs across multiple instances with AWS Batch support for **Multi-node Parallel (MNP) jobs**

Use AWS Batch to efficiently run larger-scale tightly coupled High Performance Computing (HPC) applications and distributed GPU model training without the need to launch, configure, and manage EC2 resources directly



# Typical AWS Batch Job Architecture

