



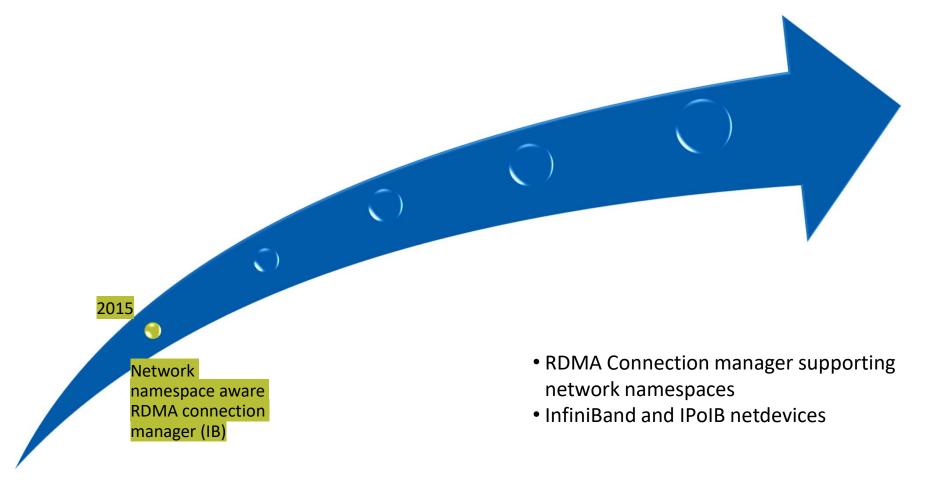
Dror Goldenberg, Parav Pandit

Mellanox Technologies June 2019



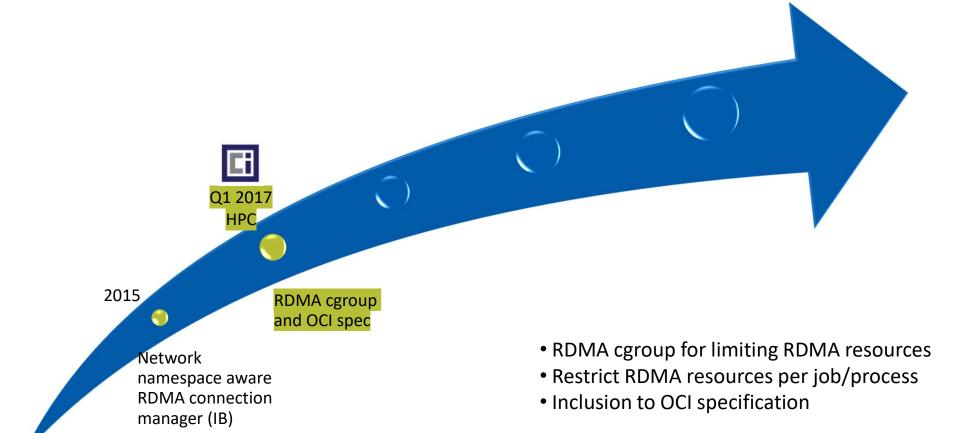
Start of RDMA container support





Extending with RDMA cgroup





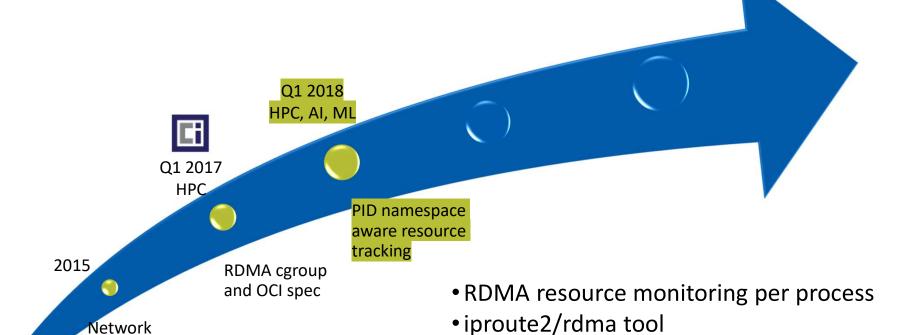
PID namespace isolation

namespace aware

RDMA connection

manager (IB)





- iproute2/rdma tool
- RDMA resources isolated per processes of a PID namespace

RoCE network namespace support

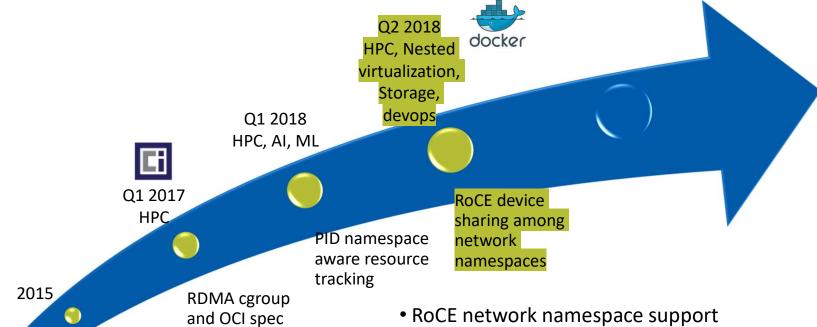
Network

namespace aware

RDMA connection

manager (IB)



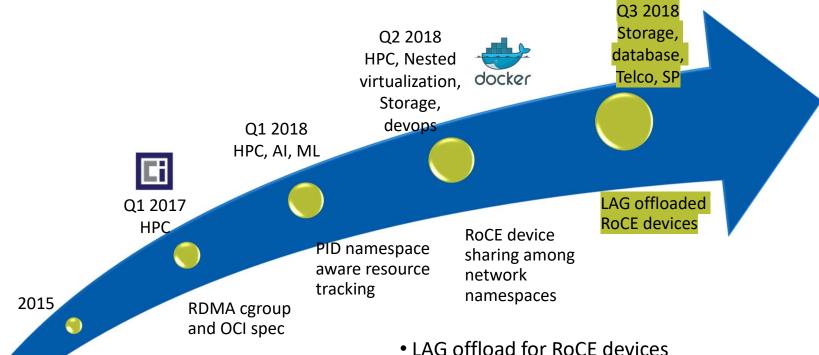


- Single RDMA device shared among multiple net namespaces
- Docker default macvlan driver, ease of deploy
- Devops, development, nested virtualization with VM, Containers use case
- Docker SR-IOV network plugin



Link/Network Resilient RoCE LAG devices





Network

• LAG officed for ROCE device

namespace aware

RDMA connection

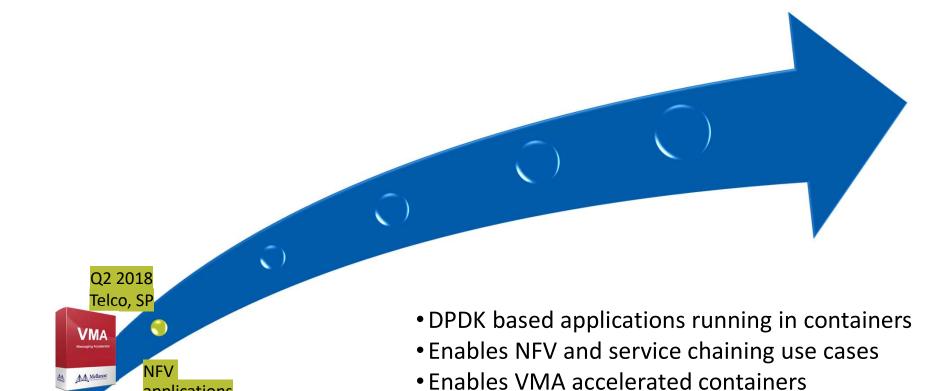
manager (IB)

- Bandwidth scaling with single RDMA and netdevice
- Transparent to container applications
- Leverages existing SR-IOV infrastructure

Enable DPDK based applications in container

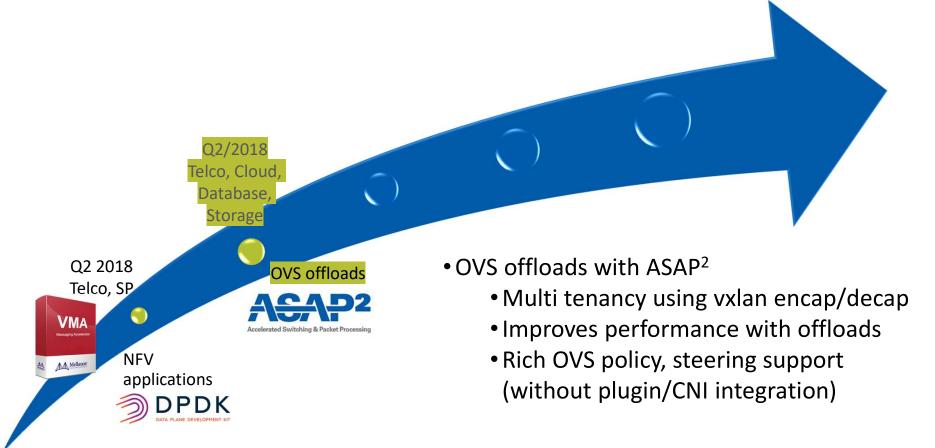
applications





OVS offloads using ASAP² technology





Kubernetes support for RDMA







Q3/2018 HPC, AI, ML, Database,



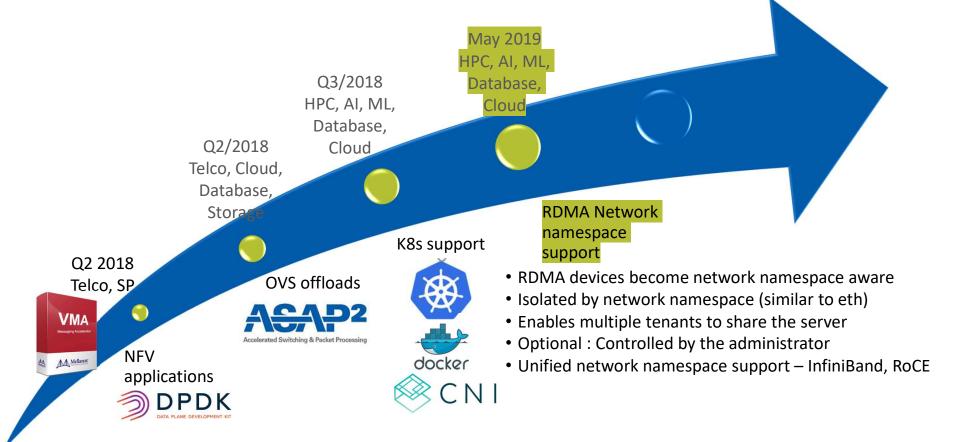




- SR-IOV CNI plugin
- SR-IOV device plugin
- Shared mode HCA device plugin

RDMA device isolation in network namespace





Openshift device plugin and CNI integration

