Missing Pieces for HPC

Container Mobility

The Production Wall

Spinning up production-like environment is…

- …not that easy

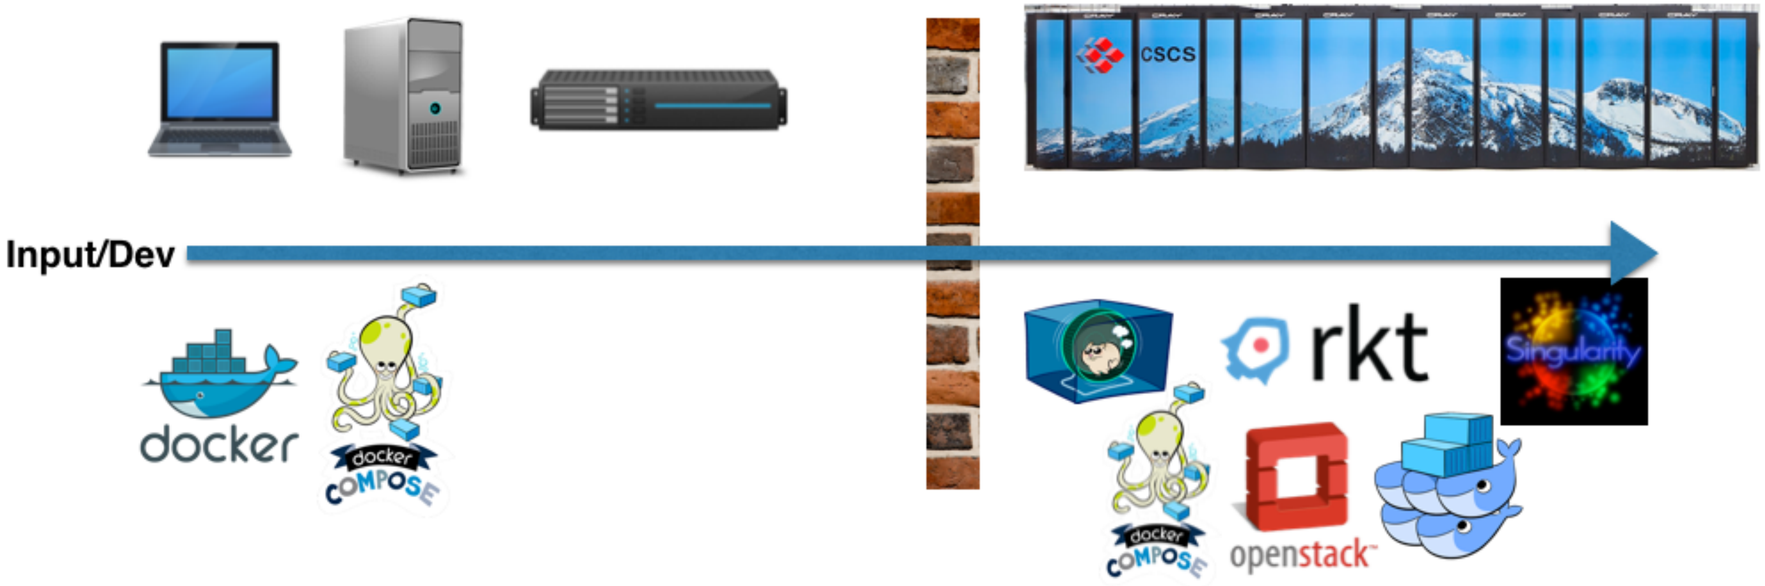

Hello Singularity

Singularity

- User-land container, leveraging some Namespaces
- creates an executable

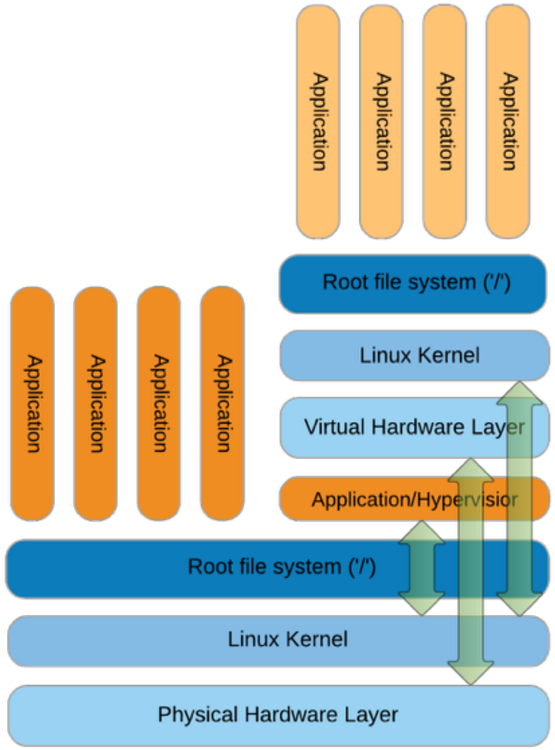

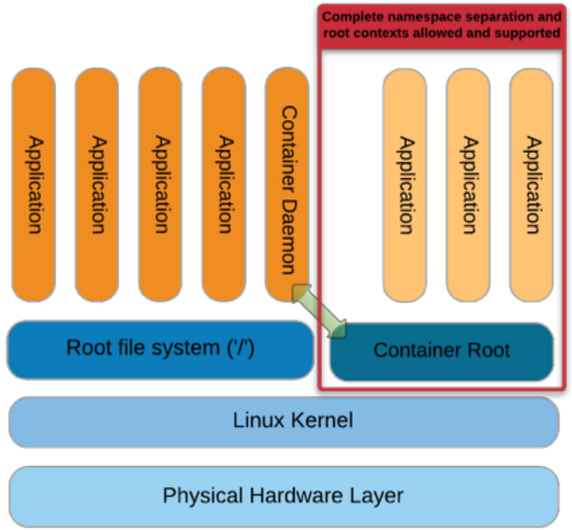

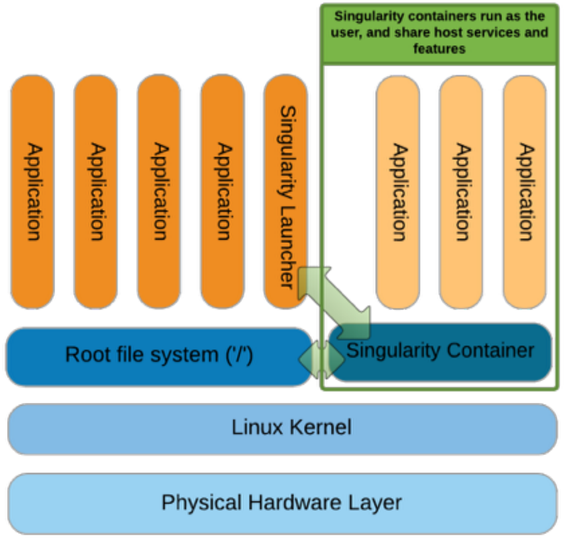

MPI Singularity

Container Tech creates new NS

- singularity start from were the user is and trims down from there
- by doing so it put a barrier on-top of what a user can do

MPI Singularity

Singularity is not bound to a daemon w/ API calls

- thus, it integrates (w/ latest) OpenMPI
- comply with the current workflow in HPC w/o changes

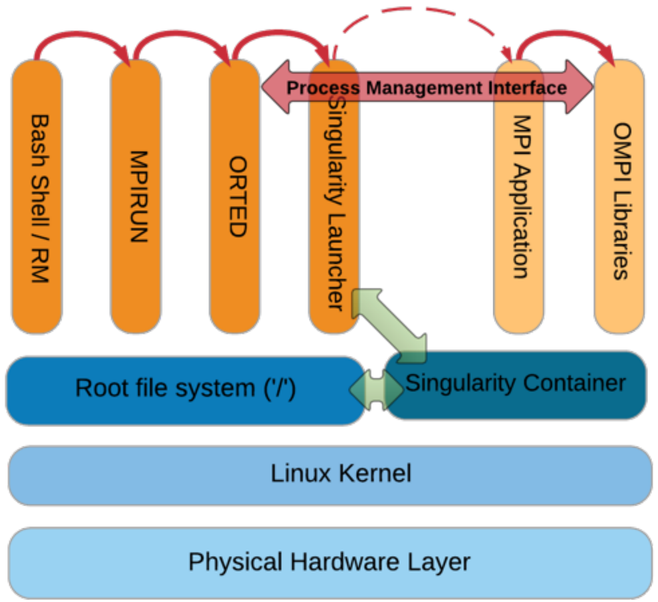


```
gmk — [screen 2: gmk@centos7-x64:~/git/ompi] — ssh gmk@gmac.dhcp.lbl....
[gmk@centos7-x64 ompi]$ time mpirun -np 4 singularity exec /tmp/Centos-7.img ./r
ing
Process 2 exiting
Process 3 exiting
Process 0 sending 10 to 1, tag 201 (4 processes in ring)
Process 0 sent to 1
Process 0 decremented value: 9
Process 0 decremented value: 8
Process 0 decremented value: 7
Process 0 decremented value: 6
Process 0 decremented value: 5
Process 0 decremented value: 4
Process 0 decremented value: 3
Process 0 decremented value: 2
Process 0 decremented value: 1
Process 0 decremented value: 0
Process 0 exiting
Process 1 exiting

real    0m0.105s
user    0m0.145s
sys     0m0.091s
[gmk@centos7-x64 ompi]$
```

MPI Singularity

Demo

RDMA Namespace

Kernel By-passing

Perfect for single tenancy

- application talks directly to hardware
- nothing in it's way

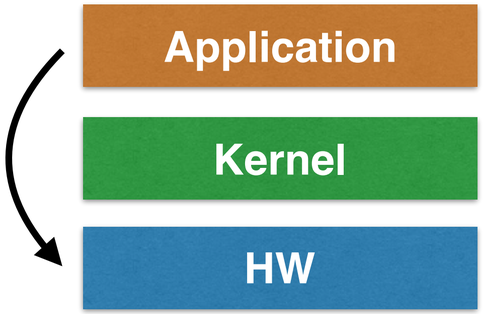

1. In VM environment we introduced SR-IOV
   - hardware is exposed as multiple devices within system
   - multiplexing is done in HW

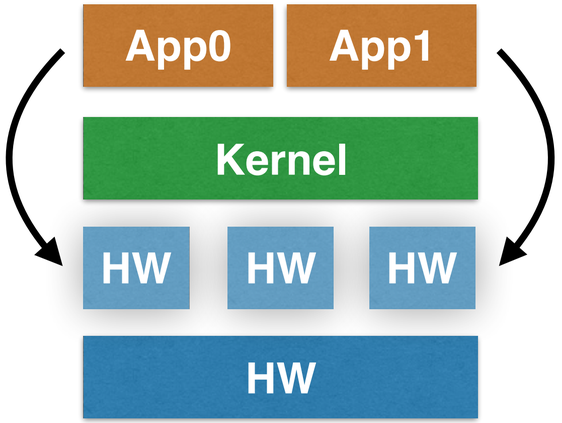

2. Linux containers are extremely volatile
   - if only we could leverage the kernel again,
     as it's aware of containers

RDMA Namespace

Mellanox pushes code for RDMA

- Namespace
- CGroups

1. By doing so, the Kernel is in charge again
   - knows about the resource needs of the all processes

RDMA Namespace #2

STATUS

- **InfiniBand RDMA CM support in v4.4**
- **RDMA cgroup patches submitted**
- **RDMA cgroup Docker patches are ready, will be submitted once kernel patches are accepted**
- **Working on RoCE net namespace support**
- **Future work**
  - InfiniBand: limit P_Key usage in verbs applications
    - Perhaps extend the RDMA cgroup
  - QoS: limit container's bandwidth usage, SL, or VLAN priority
  - Raw Ethernet support

SLURM / MPI

MPI Tricks w/ Docker

Fake a ssh-client

```
[bob@a7b1e6e98cb1 ~]$ cat /opt/qnib/src/dssh
#!/bin/bash

REMOTE_HOST=$1
shift
set -x
docker -H unix:///var/run/docker.sock exec -i -u ${USER} ${REMOTE_HOST} $@
[bob@a7b1e6e98cb1 ~]$
```

- connects via ***docker exec*** instead of ***ssh***

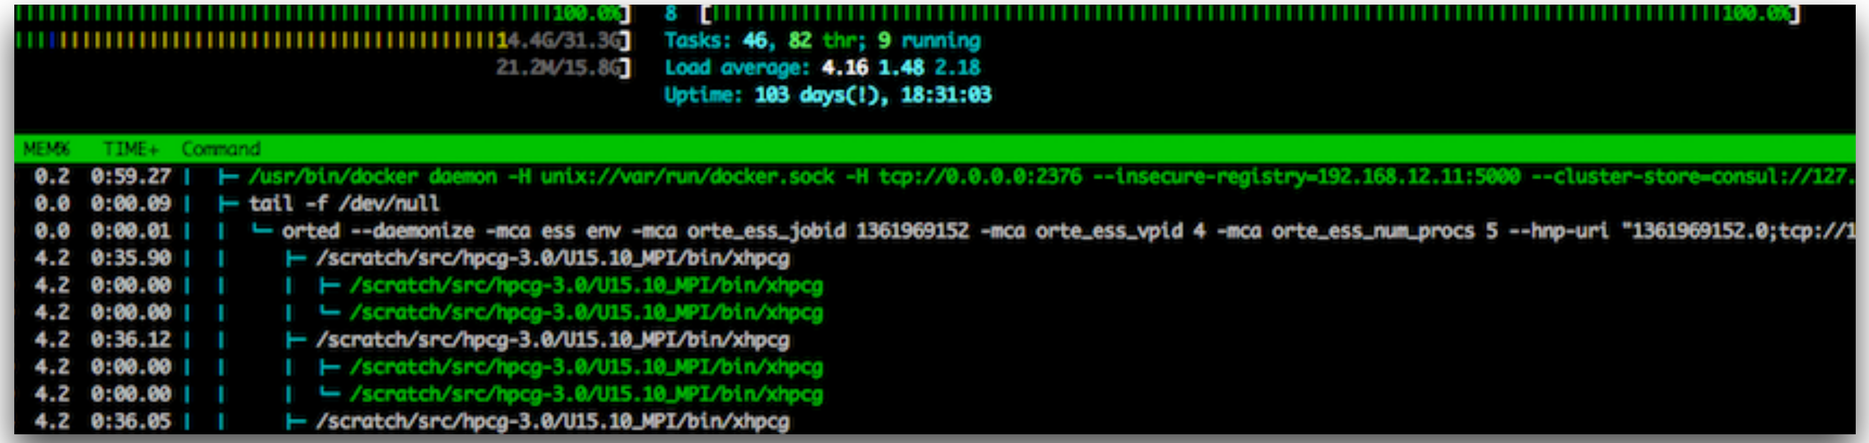

SLURM needs patching

SLURM uses ***slurmstepd*** to spawn children

```
Command
├─ /usr/sbin/slurmd
├─ /usr/sbin/munged
│  ├─ /usr/sbin/munged
│  ├─ /usr/sbin/munged
│  └─ /usr/sbin/munged
├─ slurmstepd: [752.0]
│  ├─ /usr/lib64/openmpi/bin/orted --hnp-topo-sig 0N:2S:0L3:4L2:8L1:8C:8H:x86_64 -mca orte_ess_jobid "349962240" -mca orte_ess_vpid "1" -mca orte_ess
│  │  ├─ /scratch/bin/xhpcg
│  │  │  ├─ /scratch/bin/xhpcg
│  │  │  ├─ /scratch/bin/xhpcg
│  │  │  └─ /scratch/bin/xhpcg
│  │  ├─ /scratch/bin/xhpcg
│  │  │  ├─ /scratch/bin/xhpcg
│  │  │  ├─ /scratch/bin/xhpcg
│  │  │  └─ /scratch/bin/xhpcg
│  │  ├─ /scratch/bin/xhpcg
```

If only

1. Slurmctld should use docker exec instead of fork the process directly (when in docker-mode).
2. MPI could also use docker exec to introduce process on remote system.