


# Simple custom Linux distributions with LinuxKit

Justin Cormack





# Who am I?

Engineer at Docker in Cambridge, UK.

@justincormack

# Tools for building custom Linux



# Tools for building custom Linux

Existing tools are

- very complicated
- designed for embedded systems
- very opinionated
- hard to customise
- slow to build and test

# LinuxKit



# LinuxKit

- fast to build and test
- simple configuration
- totally customisable
- runs in many different environments
- not very opinionated
- uses containers for packaging
- suitable for embedded, mainframes, supercomputers

# Started in 2015

Originally built for Docker for Mac

Needed a simple embedded, maintainable, invisible Linux

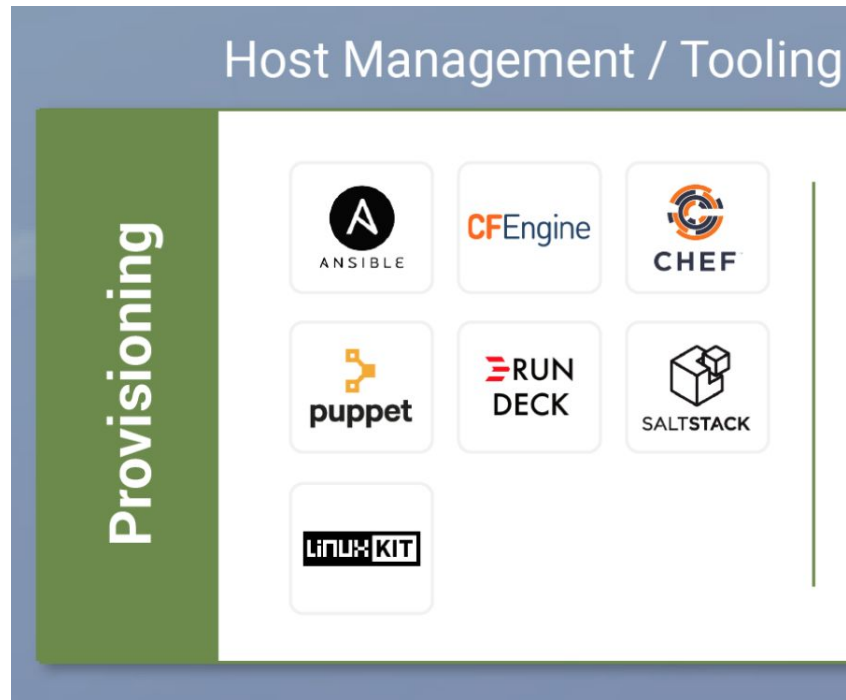
first commit: "not required: self update: treated as immutable"

This project became LinuxKit, open sourced last year



# LinuxKit is a config management tool

- defines your system configuration
- essentially just lists files
- uses containers to simplify this
- tooling to make bootable images
- built for automation



# LinuxKit

- not a Linux distro!
- it is a kit, with enough pieces to get you started
- everything can easily be replaced if required
- components specified as containers
- containers run by containerd, small container runtime
- designed to be built and tested in a CI pipeline
- build times just a minute or so
- test locally then ship to production
- minimal so boots fast
- small so secure and does not need updating so much
- Apache licensed

# Packages

- Packages are containers
- Containers are therefore the unit of defining the system
- A larger chunk than traditional systems
- Some less sharing
- Easier testability as it is a complete system
- Good service isolation
- VM isolation can be added
- Signing at the container level

# Immutable



# Sysadmins everywhere

*“As a system administrator, one of the scariest things I ever encounter is a server that’s been running for ages. If you absolutely know a system has been created via automation and never changed since the moment of creation, most of the problems disappear.”*

Chad Fowler, Trash Your Servers and Burn Your Code, 2013

# Netflix

*“In the cloud, we know exactly what we want a server to be, and if we want to change that we simply terminate it and launch a new server with a new AMI.”*

Netflix Building with Legos, 2011

# Updates



# Reprovision to update

<https://www.oreilly.com/ideas/an-introduction-to-immutable-infrastructure>

Mutable Server

Immutable Server





# State



# immutability is a name

- naming is hard
- it does not mean there is no change!
- state in traditional Unix is not very well isolated
- enforce split between code (immutable) and application data (mutable)

# Immutable does not mean stateless!

- functional programming is the analogy
  - no mutable global state
  - state change is explicit
  - state changes only made by specific parts of the code
  - aim is understandability
- state is managed and controlled, and in chosen locations
- LinuxKit has an immutable root filesystem, add writable drives for data
- controlled state mutation, not scattered all over
- persistent state changes are for data, not code or configuration

# Manage your data independently


Independent data lifecycle from code

- network storage
- cloud storage
- filesystem snapshotting (Dotmesh integration)
- replication

# LinuxKit architecture



# LinuxKit startup



sequential startup

eg network configuration, disks

services start up in parallel after initialization

***same design as pods in Kubernetes***



# Configure this from a yaml file

```
kernel:
  image: linuxkit/kernel:4.9.60
  cmdline: "console=tty0 console=ttyS0 console=ttyAMA0"
init:
  - linuxkit/init:42a92119e1ca10380e0d33e26c0cbcf85b9b3558
  - linuxkit/runc:817fdc592eac6cb7804fa1721a43a7f6e23fb50f
  - linuxkit/containerd:82be2bbb7cf83bab161ffe2a64624ba1107725ff
onboot:
  - name: dhcpd
    image: linuxkit/dhcpd:48831507404049660b960e4055f544917d90378e
    command: ["/sbin/dhcpd", "--nobackground", "-f", "/dhcpd.conf", "-1"]
services:
  - name: getty
    image: linuxkit/getty:6af22c32c98536a79230eef000e9abd06b037faa
  - name: redis
    image: redis:4.0-alpine
  capabilities:
    - CAP_NET_BIND_SERVICE
    - CAP_CHOWN
    - CAP_SETUID
    - CAP_SETGID
    - CAP_DAC_OVERRIDE
```



# note

- root filesystem is immutable
- can run from ISO, initramfs, squashfs, ...
- onboot stage is very flexible, can mount service containers from network or other arrangements if required.
- no package manager
- no possibility to update at runtime
- replace with a new image to update software
- for dynamic services can use Docker or Kubernetes on top
- removes all complexity of install, update, reboot



# Practicalities



# Simple tooling for lots of use cases

- Tooling can build most kinds of image needed to boot VMs or bare metal
  - ISO for EFI or BIOS
  - raw disk images
  - AWS AMIs
  - GCP disk format
  - QCOW2 for qemu and KVM
  - VHD
  - VMDK
  - raw kernel and initramfs
  - Raspberry Pi3 image

# Simple tooling for lots of use cases

- Simple build, push, run workflow for many common use cases
  - AWS
  - GCP
  - Azure
  - OpenStack
  - VMware Vcenter
  - Packet.net iPXE
  - Hyperkit for MacOS
  - Hyper-V for Windows
  - KVM for Linux
  - VMware Fusion
  - Virtualbox

# Simple tooling for lots of use cases

Generally (example Google Cloud)

```
linuxkit build file.yml  
linuxkit push gcp filename  
linuxkit run gcp filename
```

Some platforms have additional options. You should always use other tooling to run in production, `linuxkit run` is a development tool.

# Total control of how code runs

- Yaml config specifies the OCI configuration
- specify exactly users, capabilities, namespaces
- images can have their configuration in labels to simplify config
- modularity at the function and application level
- runs exactly what you need

# Demo



# Roadmap

- reworking build to not require Docker
  - easier to run in CI, eg in a container
  - will not require root access
  - Go code for building disk images
  - a lot of work to do here, but really useful
- more detailed application blueprints
  - a lot of work in linuxkit/kubernetes ongoing now
- remove remaining shell scripting from configuration!
  - most gone already
  - some shelling out rather than native code still
- more users and use cases!


# Summary





# Summary

- manage the configuration of a distributed system
  - use a simpler, immutable OS, just focus on the overall system
  - try our tools, they are simple but fun
  - fast tooling is way better!
  - build for automation, not for interactive use
- 
- "time to do some crazy bullshit with operating systems" Adam Jacob



THANK YOU